

project 2

● Graded

Group

Hrushik Chiluvuri

Jiyong Kwag

Nirmal Philipose Mathew

 [View or edit group](#)

Total Points

14 / 18 pts

Autograder Score

14.0 / 18.0

Failed Tests

Test the band_join (0/2)

Test the band_join with SIMD (0/2)

Passed Tests

Test the low bin search and it should always pass (2/2)

Compilation with mavx2 (2/2)

Compilation with mavx512f (2/2)

Test the low_bin_nb_arithmetic (2/2)

Test the low_bin_nb_mask (2/2)

Test the low_bin_nb_4x (2/2)

Test the low_bin_nb_simd (2/2)

Question 2

[Autograder's points is just for reference and it does not convert to the points on Carmen](#) 0 / 0 pts

✓ - 0 pts Correct

Autograder Results

Autograder Output (hidden from students)

Test band join

```
inner[0] = 0
inner[1] = 2
inner[2] = 4
inner[3] = 6
inner[4] = 8
inner[5] = 10
inner[6] = 12
inner[7] = 14
inner[8] = 16
inner[9] = 18
inner[10] = 20
inner[11] = 22
inner[12] = 24
inner[13] = 26
inner[14] = 28
inner[15] = 30
inner[16] = 32
inner[17] = 34
inner[18] = 36
inner[19] = 38
inner[20] = 40
inner[21] = 42
inner[22] = 44
inner[23] = 46
inner[24] = 48
inner[25] = 50
inner[26] = 52
inner[27] = 54
inner[28] = 56
inner[29] = 58
inner[30] = 60
inner[31] = 62
inner[32] = 64
inner[33] = 66
inner[34] = 68
inner[35] = 70
inner[36] = 72
inner[37] = 74
inner[38] = 76
inner[39] = 78
inner[40] = 80
inner[41] = 82
inner[42] = 84
inner[43] = 86
inner[44] = 88
inner[45] = 90
inner[46] = 92
```

inner[47] = 94
inner[48] = 96
inner[49] = 98
inner[50] = 100
inner[51] = 102
inner[52] = 104
inner[53] = 106
inner[54] = 108
inner[55] = 110
inner[56] = 112
inner[57] = 114
inner[58] = 116
inner[59] = 118
inner[60] = 120
inner[61] = 122
inner[62] = 124
inner[63] = 126
inner[64] = 128
inner[65] = 130
inner[66] = 132
inner[67] = 134
inner[68] = 136
inner[69] = 138
inner[70] = 140
inner[71] = 142
inner[72] = 144
inner[73] = 146
inner[74] = 148
inner[75] = 150
inner[76] = 152
inner[77] = 154
inner[78] = 156
inner[79] = 158
inner[80] = 160
inner[81] = 162
inner[82] = 164
inner[83] = 166
inner[84] = 168
inner[85] = 170
inner[86] = 172
inner[87] = 174
inner[88] = 176
inner[89] = 178
inner[90] = 180
inner[91] = 182
inner[92] = 184
inner[93] = 186
inner[94] = 188
inner[95] = 190

```
inner[96] = 192
inner[97] = 194
inner[98] = 196
inner[99] = 198
outer[0] = 80
outer[1] = 78
outer[2] = 76
outer[3] = 74
outer[4] = 72
outer[5] = 70
outer[6] = 68
outer[7] = 66
outer[8] = 64
outer[9] = 62
outer[10] = 60
outer[11] = 58
outer[12] = 56
outer[13] = 54
outer[14] = 52
outer[15] = 50
outer[16] = 48
outer[17] = 46
outer[18] = 44
outer[19] = 42
outer[20] = 40
outer[21] = 38
outer[22] = 36
outer[23] = 34
outer[24] = 32
outer[25] = 30
outer[26] = 28
outer[27] = 26
outer[28] = 24
outer[29] = 22
outer[30] = 20
outer[31] = 18
outer[32] = 16
outer[33] = 14
outer[34] = 12
outer[35] = 10
outer[36] = 8
outer[37] = 6
outer[38] = 4
```

band join with band 4 return 156 tuples

The return number 156 is incorrect

(0, 38)

(1, 37)

(1, 38)

(2, 36)
(2, 37)
(2, 38)
(3, 35)
(3, 36)
(3, 37)
(3, 38)
(4, 34)
(4, 35)
(4, 36)
(4, 37)
(5, 33)
(5, 34)
(5, 35)
(5, 36)
(6, 32)
(6, 33)
(6, 34)
(6, 35)
(7, 31)
(7, 32)
(7, 33)
(7, 34)
(8, 30)
(8, 31)
(8, 32)
(8, 33)
(9, 29)
(9, 30)
(9, 31)
(9, 32)
(10, 28)
(10, 29)
(10, 30)
(10, 31)
(11, 27)
(11, 28)
(11, 29)
(11, 30)
(12, 26)
(12, 27)
(12, 28)
(12, 29)
(13, 25)
(13, 26)
(13, 27)
(13, 28)
(14, 24)
(14, 25)

(14, 26)
(14, 27)
(15, 23)
(15, 24)
(15, 25)
(15, 26)
(16, 22)
(16, 23)
(16, 24)
(16, 25)
(17, 21)
(17, 22)
(17, 23)
(17, 24)
(18, 20)
(18, 21)
(18, 22)
(18, 23)
(19, 19)
(19, 20)
(19, 21)
(19, 22)
(20, 18)
(20, 19)
(20, 20)
(20, 21)
(21, 17)
(21, 18)
(21, 19)
(21, 20)
(22, 16)
(22, 17)
(22, 18)
(22, 19)
(23, 15)
(23, 16)
(23, 17)
(23, 18)
(24, 14)
(24, 15)
(24, 16)
(24, 17)
(25, 13)
(25, 14)
(25, 15)
(25, 16)
(26, 12)
(26, 13)
(26, 14)

(26, 15)
(27, 11)
(27, 12)
(27, 13)
(27, 14)
(28, 10)
(28, 11)
(28, 12)
(28, 13)
(29, 9)
(29, 10)
(29, 11)
(29, 12)
(30, 8)
(30, 9)
(30, 10)
(30, 11)
(31, 7)
(31, 8)
(31, 9)
(31, 10)
(32, 6)
(32, 7)
(32, 8)
(32, 9)
(33, 5)
(33, 6)
(33, 7)
(33, 8)
(34, 4)
(34, 5)
(34, 6)
(34, 7)
(35, 3)
(35, 4)
(35, 5)
(35, 6)
(36, 2)
(36, 3)
(36, 4)
(36, 5)
(37, 1)
(37, 2)
(37, 3)
(37, 4)
(38, 0)
(38, 1)
(38, 2)
(38, 3)

(39, 0)

(39, 1)

(39, 2)

(40, 0)

(40, 1)

(41, 0)

Some error in the band join, correct result should be

(0, 38)

(1, 37)

(1, 38)

(2, 36)

(2, 37)

(2, 38)

(3, 35)

(3, 36)

(3, 37)

(3, 38)

(4, 34)

(4, 35)

(4, 36)

(4, 37)

(4, 38)

(5, 33)

(5, 34)

(5, 35)

(5, 36)

(5, 37)

(6, 32)

(6, 33)

(6, 34)

(6, 35)

(6, 36)

(7, 31)

(7, 32)

(7, 33)

(7, 34)

(7, 35)

(8, 30)

(8, 31)

(8, 32)

(8, 33)

(8, 34)

(9, 29)

(9, 30)

(9, 31)

(9, 32)

(9, 33)

(10, 28)

(10, 29)

(10, 30)
(10, 31)
(10, 32)
(11, 27)
(11, 28)
(11, 29)
(11, 30)
(11, 31)
(12, 26)
(12, 27)
(12, 28)
(12, 29)
(12, 30)
(13, 25)
(13, 26)
(13, 27)
(13, 28)
(13, 29)
(14, 24)
(14, 25)
(14, 26)
(14, 27)
(14, 28)
(15, 23)
(15, 24)
(15, 25)
(15, 26)
(15, 27)
(16, 22)
(16, 23)
(16, 24)
(16, 25)
(16, 26)
(17, 21)
(17, 22)
(17, 23)
(17, 24)
(17, 25)
(18, 20)
(18, 21)
(18, 22)
(18, 23)
(18, 24)
(19, 19)
(19, 20)
(19, 21)
(19, 22)
(19, 23)
(20, 18)

(20, 19)
(20, 20)
(20, 21)
(20, 22)
(21, 17)
(21, 18)
(21, 19)
(21, 20)
(21, 21)
(22, 16)
(22, 17)
(22, 18)
(22, 19)
(22, 20)
(23, 15)
(23, 16)
(23, 17)
(23, 18)
(23, 19)
(24, 14)
(24, 15)
(24, 16)
(24, 17)
(24, 18)
(25, 13)
(25, 14)
(25, 15)
(25, 16)
(25, 17)
(26, 12)
(26, 13)
(26, 14)
(26, 15)
(26, 16)
(27, 11)
(27, 12)
(27, 13)
(27, 14)
(27, 15)
(28, 10)
(28, 11)
(28, 12)
(28, 13)
(28, 14)
(29, 9)
(29, 10)
(29, 11)
(29, 12)
(29, 13)

(30, 8)
(30, 9)
(30, 10)
(30, 11)
(30, 12)
(31, 7)
(31, 8)
(31, 9)
(31, 10)
(31, 11)
(32, 6)
(32, 7)
(32, 8)
(32, 9)
(32, 10)
(33, 5)
(33, 6)
(33, 7)
(33, 8)
(33, 9)
(34, 4)
(34, 5)
(34, 6)
(34, 7)
(34, 8)
(35, 3)
(35, 4)
(35, 5)
(35, 6)
(35, 7)
(36, 2)
(36, 3)
(36, 4)
(36, 5)
(36, 6)
(37, 1)
(37, 2)
(37, 3)
(37, 4)
(37, 5)
(38, 0)
(38, 1)
(38, 2)
(38, 3)
(38, 4)
(39, 0)
(39, 1)
(39, 2)
(39, 3)

(40, 0)

(40, 1)

(40, 2)

(41, 0)

(41, 1)

(42, 0)

Some error happens, will try with a reverse order

inner[0] = 198

inner[1] = 196

inner[2] = 194

inner[3] = 192

inner[4] = 190

inner[5] = 188

inner[6] = 186

inner[7] = 184

inner[8] = 182

inner[9] = 180

inner[10] = 178

inner[11] = 176

inner[12] = 174

inner[13] = 172

inner[14] = 170

inner[15] = 168

inner[16] = 166

inner[17] = 164

inner[18] = 162

inner[19] = 160

inner[20] = 158

inner[21] = 156

inner[22] = 154

inner[23] = 152

inner[24] = 150

inner[25] = 148

inner[26] = 146

inner[27] = 144

inner[28] = 142

inner[29] = 140

inner[30] = 138

inner[31] = 136

inner[32] = 134

inner[33] = 132

inner[34] = 130

inner[35] = 128

inner[36] = 126

inner[37] = 124

inner[38] = 122

inner[39] = 120

inner[40] = 118

inner[41] = 116

inner[42] = 114
inner[43] = 112
inner[44] = 110
inner[45] = 108
inner[46] = 106
inner[47] = 104
inner[48] = 102
inner[49] = 100
inner[50] = 98
inner[51] = 96
inner[52] = 94
inner[53] = 92
inner[54] = 90
inner[55] = 88
inner[56] = 86
inner[57] = 84
inner[58] = 82
inner[59] = 80
inner[60] = 78
inner[61] = 76
inner[62] = 74
inner[63] = 72
inner[64] = 70
inner[65] = 68
inner[66] = 66
inner[67] = 64
inner[68] = 62
inner[69] = 60
inner[70] = 58
inner[71] = 56
inner[72] = 54
inner[73] = 52
inner[74] = 50
inner[75] = 48
inner[76] = 46
inner[77] = 44
inner[78] = 42
inner[79] = 40
inner[80] = 38
inner[81] = 36
inner[82] = 34
inner[83] = 32
inner[84] = 30
inner[85] = 28
inner[86] = 26
inner[87] = 24
inner[88] = 22
inner[89] = 20
inner[90] = 18

```
inner[91] = 16
inner[92] = 14
inner[93] = 12
inner[94] = 10
inner[95] = 8
inner[96] = 6
inner[97] = 4
inner[98] = 2
inner[99] = 0
outer[0] = 4
outer[1] = 6
outer[2] = 8
outer[3] = 10
outer[4] = 12
outer[5] = 14
outer[6] = 16
outer[7] = 18
outer[8] = 20
outer[9] = 22
outer[10] = 24
outer[11] = 26
outer[12] = 28
outer[13] = 30
outer[14] = 32
outer[15] = 34
outer[16] = 36
outer[17] = 38
outer[18] = 40
outer[19] = 42
outer[20] = 44
outer[21] = 46
outer[22] = 48
outer[23] = 50
outer[24] = 52
outer[25] = 54
outer[26] = 56
outer[27] = 58
outer[28] = 60
outer[29] = 62
outer[30] = 64
outer[31] = 66
outer[32] = 68
outer[33] = 70
outer[34] = 72
outer[35] = 74
outer[36] = 76
outer[37] = 78
outer[38] = 80
```

band join with band 4 return 0 tuples

The return number 0 is incorrect

Some error in the band join, correct result should be

(57, 38)

(58, 37)

(58, 38)

(59, 36)

(59, 37)

(59, 38)

(60, 35)

(60, 36)

(60, 37)

(60, 38)

(61, 34)

(61, 35)

(61, 36)

(61, 37)

(61, 38)

(62, 33)

(62, 34)

(62, 35)

(62, 36)

(62, 37)

(63, 32)

(63, 33)

(63, 34)

(63, 35)

(63, 36)

(64, 31)

(64, 32)

(64, 33)

(64, 34)

(64, 35)

(65, 30)

(65, 31)

(65, 32)

(65, 33)

(65, 34)

(66, 29)

(66, 30)

(66, 31)

(66, 32)

(66, 33)

(67, 28)

(67, 29)

(67, 30)

(67, 31)

(67, 32)

(68, 27)

(68, 28)
(68, 29)
(68, 30)
(68, 31)
(69, 26)
(69, 27)
(69, 28)
(69, 29)
(69, 30)
(70, 25)
(70, 26)
(70, 27)
(70, 28)
(70, 29)
(71, 24)
(71, 25)
(71, 26)
(71, 27)
(71, 28)
(72, 23)
(72, 24)
(72, 25)
(72, 26)
(72, 27)
(73, 22)
(73, 23)
(73, 24)
(73, 25)
(73, 26)
(74, 21)
(74, 22)
(74, 23)
(74, 24)
(74, 25)
(75, 20)
(75, 21)
(75, 22)
(75, 23)
(75, 24)
(76, 19)
(76, 20)
(76, 21)
(76, 22)
(76, 23)
(77, 18)
(77, 19)
(77, 20)
(77, 21)
(77, 22)

(78, 17)
(78, 18)
(78, 19)
(78, 20)
(78, 21)
(79, 16)
(79, 17)
(79, 18)
(79, 19)
(79, 20)
(80, 15)
(80, 16)
(80, 17)
(80, 18)
(80, 19)
(81, 14)
(81, 15)
(81, 16)
(81, 17)
(81, 18)
(82, 13)
(82, 14)
(82, 15)
(82, 16)
(82, 17)
(83, 12)
(83, 13)
(83, 14)
(83, 15)
(83, 16)
(84, 11)
(84, 12)
(84, 13)
(84, 14)
(84, 15)
(85, 10)
(85, 11)
(85, 12)
(85, 13)
(85, 14)
(86, 9)
(86, 10)
(86, 11)
(86, 12)
(86, 13)
(87, 8)
(87, 9)
(87, 10)
(87, 11)

(87, 12)
(88, 7)
(88, 8)
(88, 9)
(88, 10)
(88, 11)
(89, 6)
(89, 7)
(89, 8)
(89, 9)
(89, 10)
(90, 5)
(90, 6)
(90, 7)
(90, 8)
(90, 9)
(91, 4)
(91, 5)
(91, 6)
(91, 7)
(91, 8)
(92, 3)
(92, 4)
(92, 5)
(92, 6)
(92, 7)
(93, 2)
(93, 3)
(93, 4)
(93, 5)
(93, 6)
(94, 1)
(94, 2)
(94, 3)
(94, 4)
(94, 5)
(95, 0)
(95, 1)
(95, 2)
(95, 3)
(95, 4)
(96, 0)
(96, 1)
(96, 2)
(96, 3)
(97, 0)
(97, 1)
(97, 2)
(98, 0)

(98, 1)

(99, 0)

Test band join with SIMD

inner[0] = 0

inner[1] = 2

inner[2] = 4

inner[3] = 6

inner[4] = 8

inner[5] = 10

inner[6] = 12

inner[7] = 14

inner[8] = 16

inner[9] = 18

inner[10] = 20

inner[11] = 22

inner[12] = 24

inner[13] = 26

inner[14] = 28

inner[15] = 30

inner[16] = 32

inner[17] = 34

inner[18] = 36

inner[19] = 38

inner[20] = 40

inner[21] = 42

inner[22] = 44

inner[23] = 46

inner[24] = 48

inner[25] = 50

inner[26] = 52

inner[27] = 54

inner[28] = 56

inner[29] = 58

inner[30] = 60

inner[31] = 62

inner[32] = 64

inner[33] = 66

inner[34] = 68

inner[35] = 70

inner[36] = 72

inner[37] = 74

inner[38] = 76

inner[39] = 78

inner[40] = 80

inner[41] = 82

inner[42] = 84

inner[43] = 86

inner[44] = 88

inner[45] = 90

inner[46] = 92
inner[47] = 94
inner[48] = 96
inner[49] = 98
inner[50] = 100
inner[51] = 102
inner[52] = 104
inner[53] = 106
inner[54] = 108
inner[55] = 110
inner[56] = 112
inner[57] = 114
inner[58] = 116
inner[59] = 118
inner[60] = 120
inner[61] = 122
inner[62] = 124
inner[63] = 126
inner[64] = 128
inner[65] = 130
inner[66] = 132
inner[67] = 134
inner[68] = 136
inner[69] = 138
inner[70] = 140
inner[71] = 142
inner[72] = 144
inner[73] = 146
inner[74] = 148
inner[75] = 150
inner[76] = 152
inner[77] = 154
inner[78] = 156
inner[79] = 158
inner[80] = 160
inner[81] = 162
inner[82] = 164
inner[83] = 166
inner[84] = 168
inner[85] = 170
inner[86] = 172
inner[87] = 174
inner[88] = 176
inner[89] = 178
inner[90] = 180
inner[91] = 182
inner[92] = 184
inner[93] = 186
inner[94] = 188

```
inner[95] = 190
inner[96] = 192
inner[97] = 194
inner[98] = 196
inner[99] = 198
outer[0] = 80
outer[1] = 78
outer[2] = 76
outer[3] = 74
outer[4] = 72
outer[5] = 70
outer[6] = 68
outer[7] = 66
outer[8] = 64
outer[9] = 62
outer[10] = 60
outer[11] = 58
outer[12] = 56
outer[13] = 54
outer[14] = 52
outer[15] = 50
outer[16] = 48
outer[17] = 46
outer[18] = 44
outer[19] = 42
outer[20] = 40
outer[21] = 38
outer[22] = 36
outer[23] = 34
outer[24] = 32
outer[25] = 30
outer[26] = 28
outer[27] = 26
outer[28] = 24
outer[29] = 22
outer[30] = 20
outer[31] = 18
outer[32] = 16
outer[33] = 14
outer[34] = 12
outer[35] = 10
outer[36] = 8
outer[37] = 6
outer[38] = 4
band join with band 4 return 156 tuples
The return number 156 is incorrect
( 0, 38)
( 1, 37)
( 1, 38)
```

(2, 36)
(2, 37)
(2, 38)
(3, 35)
(3, 36)
(3, 37)
(3, 38)
(4, 34)
(4, 35)
(4, 36)
(4, 37)
(5, 33)
(5, 34)
(5, 35)
(5, 36)
(6, 32)
(6, 33)
(6, 34)
(6, 35)
(7, 31)
(7, 32)
(7, 33)
(7, 34)
(8, 30)
(8, 31)
(8, 32)
(8, 33)
(9, 29)
(9, 30)
(9, 31)
(9, 32)
(10, 28)
(10, 29)
(10, 30)
(10, 31)
(11, 27)
(11, 28)
(11, 29)
(11, 30)
(12, 26)
(12, 27)
(12, 28)
(12, 29)
(13, 25)
(13, 26)
(13, 27)
(13, 28)
(14, 24)
(14, 25)

(14, 26)
(14, 27)
(15, 23)
(15, 24)
(15, 25)
(15, 26)
(16, 22)
(16, 23)
(16, 24)
(16, 25)
(17, 21)
(17, 22)
(17, 23)
(17, 24)
(18, 20)
(18, 21)
(18, 22)
(18, 23)
(19, 19)
(19, 20)
(19, 21)
(19, 22)
(20, 18)
(20, 19)
(20, 20)
(20, 21)
(21, 17)
(21, 18)
(21, 19)
(21, 20)
(22, 16)
(22, 17)
(22, 18)
(22, 19)
(23, 15)
(23, 16)
(23, 17)
(23, 18)
(24, 14)
(24, 15)
(24, 16)
(24, 17)
(25, 13)
(25, 14)
(25, 15)
(25, 16)
(26, 12)
(26, 13)
(26, 14)

(26, 15)
(27, 11)
(27, 12)
(27, 13)
(27, 14)
(28, 10)
(28, 11)
(28, 12)
(28, 13)
(29, 9)
(29, 10)
(29, 11)
(29, 12)
(30, 8)
(30, 9)
(30, 10)
(30, 11)
(31, 7)
(31, 8)
(31, 9)
(31, 10)
(32, 6)
(32, 7)
(32, 8)
(32, 9)
(33, 5)
(33, 6)
(33, 7)
(33, 8)
(34, 4)
(34, 5)
(34, 6)
(34, 7)
(35, 3)
(35, 4)
(35, 5)
(35, 6)
(36, 2)
(36, 3)
(36, 4)
(36, 5)
(37, 1)
(37, 2)
(37, 3)
(37, 4)
(38, 0)
(38, 1)
(38, 2)
(38, 3)

(39, 0)

(39, 1)

(39, 2)

(40, 0)

(40, 1)

(41, 0)

Some error in the band join, correct result should be

(0, 38)

(1, 37)

(1, 38)

(2, 36)

(2, 37)

(2, 38)

(3, 35)

(3, 36)

(3, 37)

(3, 38)

(4, 34)

(4, 35)

(4, 36)

(4, 37)

(4, 38)

(5, 33)

(5, 34)

(5, 35)

(5, 36)

(5, 37)

(6, 32)

(6, 33)

(6, 34)

(6, 35)

(6, 36)

(7, 31)

(7, 32)

(7, 33)

(7, 34)

(7, 35)

(8, 30)

(8, 31)

(8, 32)

(8, 33)

(8, 34)

(9, 29)

(9, 30)

(9, 31)

(9, 32)

(9, 33)

(10, 28)

(10, 29)

(10, 30)
(10, 31)
(10, 32)
(11, 27)
(11, 28)
(11, 29)
(11, 30)
(11, 31)
(12, 26)
(12, 27)
(12, 28)
(12, 29)
(12, 30)
(13, 25)
(13, 26)
(13, 27)
(13, 28)
(13, 29)
(14, 24)
(14, 25)
(14, 26)
(14, 27)
(14, 28)
(15, 23)
(15, 24)
(15, 25)
(15, 26)
(15, 27)
(16, 22)
(16, 23)
(16, 24)
(16, 25)
(16, 26)
(17, 21)
(17, 22)
(17, 23)
(17, 24)
(17, 25)
(18, 20)
(18, 21)
(18, 22)
(18, 23)
(18, 24)
(19, 19)
(19, 20)
(19, 21)
(19, 22)
(19, 23)
(20, 18)

(20, 19)
(20, 20)
(20, 21)
(20, 22)
(21, 17)
(21, 18)
(21, 19)
(21, 20)
(21, 21)
(22, 16)
(22, 17)
(22, 18)
(22, 19)
(22, 20)
(23, 15)
(23, 16)
(23, 17)
(23, 18)
(23, 19)
(24, 14)
(24, 15)
(24, 16)
(24, 17)
(24, 18)
(25, 13)
(25, 14)
(25, 15)
(25, 16)
(25, 17)
(26, 12)
(26, 13)
(26, 14)
(26, 15)
(26, 16)
(27, 11)
(27, 12)
(27, 13)
(27, 14)
(27, 15)
(28, 10)
(28, 11)
(28, 12)
(28, 13)
(28, 14)
(29, 9)
(29, 10)
(29, 11)
(29, 12)
(29, 13)

(30, 8)
(30, 9)
(30, 10)
(30, 11)
(30, 12)
(31, 7)
(31, 8)
(31, 9)
(31, 10)
(31, 11)
(32, 6)
(32, 7)
(32, 8)
(32, 9)
(32, 10)
(33, 5)
(33, 6)
(33, 7)
(33, 8)
(33, 9)
(34, 4)
(34, 5)
(34, 6)
(34, 7)
(34, 8)
(35, 3)
(35, 4)
(35, 5)
(35, 6)
(35, 7)
(36, 2)
(36, 3)
(36, 4)
(36, 5)
(36, 6)
(37, 1)
(37, 2)
(37, 3)
(37, 4)
(37, 5)
(38, 0)
(38, 1)
(38, 2)
(38, 3)
(38, 4)
(39, 0)
(39, 1)
(39, 2)
(39, 3)

(40, 0)

(40, 1)

(40, 2)

(41, 0)

(41, 1)

(42, 0)

Some error happens, will try with a reverse order

inner[0] = 198

inner[1] = 196

inner[2] = 194

inner[3] = 192

inner[4] = 190

inner[5] = 188

inner[6] = 186

inner[7] = 184

inner[8] = 182

inner[9] = 180

inner[10] = 178

inner[11] = 176

inner[12] = 174

inner[13] = 172

inner[14] = 170

inner[15] = 168

inner[16] = 166

inner[17] = 164

inner[18] = 162

inner[19] = 160

inner[20] = 158

inner[21] = 156

inner[22] = 154

inner[23] = 152

inner[24] = 150

inner[25] = 148

inner[26] = 146

inner[27] = 144

inner[28] = 142

inner[29] = 140

inner[30] = 138

inner[31] = 136

inner[32] = 134

inner[33] = 132

inner[34] = 130

inner[35] = 128

inner[36] = 126

inner[37] = 124

inner[38] = 122

inner[39] = 120

inner[40] = 118

inner[41] = 116

inner[42] = 114
inner[43] = 112
inner[44] = 110
inner[45] = 108
inner[46] = 106
inner[47] = 104
inner[48] = 102
inner[49] = 100
inner[50] = 98
inner[51] = 96
inner[52] = 94
inner[53] = 92
inner[54] = 90
inner[55] = 88
inner[56] = 86
inner[57] = 84
inner[58] = 82
inner[59] = 80
inner[60] = 78
inner[61] = 76
inner[62] = 74
inner[63] = 72
inner[64] = 70
inner[65] = 68
inner[66] = 66
inner[67] = 64
inner[68] = 62
inner[69] = 60
inner[70] = 58
inner[71] = 56
inner[72] = 54
inner[73] = 52
inner[74] = 50
inner[75] = 48
inner[76] = 46
inner[77] = 44
inner[78] = 42
inner[79] = 40
inner[80] = 38
inner[81] = 36
inner[82] = 34
inner[83] = 32
inner[84] = 30
inner[85] = 28
inner[86] = 26
inner[87] = 24
inner[88] = 22
inner[89] = 20
inner[90] = 18

```
inner[91] = 16
inner[92] = 14
inner[93] = 12
inner[94] = 10
inner[95] = 8
inner[96] = 6
inner[97] = 4
inner[98] = 2
inner[99] = 0
outer[0] = 4
outer[1] = 6
outer[2] = 8
outer[3] = 10
outer[4] = 12
outer[5] = 14
outer[6] = 16
outer[7] = 18
outer[8] = 20
outer[9] = 22
outer[10] = 24
outer[11] = 26
outer[12] = 28
outer[13] = 30
outer[14] = 32
outer[15] = 34
outer[16] = 36
outer[17] = 38
outer[18] = 40
outer[19] = 42
outer[20] = 44
outer[21] = 46
outer[22] = 48
outer[23] = 50
outer[24] = 52
outer[25] = 54
outer[26] = 56
outer[27] = 58
outer[28] = 60
outer[29] = 62
outer[30] = 64
outer[31] = 66
outer[32] = 68
outer[33] = 70
outer[34] = 72
outer[35] = 74
outer[36] = 76
outer[37] = 78
outer[38] = 80
band join with band 4 return 0 tuples
```


The return number 0 is incorrect

Some error in the band join, correct result should be

(57, 38)
(58, 37)
(58, 38)
(59, 36)
(59, 37)
(59, 38)
(60, 35)
(60, 36)
(60, 37)
(60, 38)
(61, 34)
(61, 35)
(61, 36)
(61, 37)
(61, 38)
(62, 33)
(62, 34)
(62, 35)
(62, 36)
(62, 37)
(63, 32)
(63, 33)
(63, 34)
(63, 35)
(63, 36)
(64, 31)
(64, 32)
(64, 33)
(64, 34)
(64, 35)
(65, 30)
(65, 31)
(65, 32)
(65, 33)
(65, 34)
(66, 29)
(66, 30)
(66, 31)
(66, 32)
(66, 33)
(67, 28)
(67, 29)
(67, 30)
(67, 31)
(67, 32)
(68, 27)
(68, 28)

(68, 29)
(68, 30)
(68, 31)
(69, 26)
(69, 27)
(69, 28)
(69, 29)
(69, 30)
(70, 25)
(70, 26)
(70, 27)
(70, 28)
(70, 29)
(71, 24)
(71, 25)
(71, 26)
(71, 27)
(71, 28)
(72, 23)
(72, 24)
(72, 25)
(72, 26)
(72, 27)
(73, 22)
(73, 23)
(73, 24)
(73, 25)
(73, 26)
(74, 21)
(74, 22)
(74, 23)
(74, 24)
(74, 25)
(75, 20)
(75, 21)
(75, 22)
(75, 23)
(75, 24)
(76, 19)
(76, 20)
(76, 21)
(76, 22)
(76, 23)
(77, 18)
(77, 19)
(77, 20)
(77, 21)
(77, 22)
(78, 17)

(78, 18)
(78, 19)
(78, 20)
(78, 21)
(79, 16)
(79, 17)
(79, 18)
(79, 19)
(79, 20)
(80, 15)
(80, 16)
(80, 17)
(80, 18)
(80, 19)
(81, 14)
(81, 15)
(81, 16)
(81, 17)
(81, 18)
(82, 13)
(82, 14)
(82, 15)
(82, 16)
(82, 17)
(83, 12)
(83, 13)
(83, 14)
(83, 15)
(83, 16)
(84, 11)
(84, 12)
(84, 13)
(84, 14)
(84, 15)
(85, 10)
(85, 11)
(85, 12)
(85, 13)
(85, 14)
(86, 9)
(86, 10)
(86, 11)
(86, 12)
(86, 13)
(87, 8)
(87, 9)
(87, 10)
(87, 11)
(87, 12)

(88, 7)
(88, 8)
(88, 9)
(88, 10)
(88, 11)
(89, 6)
(89, 7)
(89, 8)
(89, 9)
(89, 10)
(90, 5)
(90, 6)
(90, 7)
(90, 8)
(90, 9)
(91, 4)
(91, 5)
(91, 6)
(91, 7)
(91, 8)
(92, 3)
(92, 4)
(92, 5)
(92, 6)
(92, 7)
(93, 2)
(93, 3)
(93, 4)
(93, 5)
(93, 6)
(94, 1)
(94, 2)
(94, 3)
(94, 4)
(94, 5)
(95, 0)
(95, 1)
(95, 2)
(95, 3)
(95, 4)
(96, 0)
(96, 1)
(96, 2)
(96, 3)
(97, 0)
(97, 1)
(97, 2)
(98, 0)
(98, 1)

(99, 0)

Test low_bin_search

data[0] = 0

data[1] = 2

data[2] = 4

data[3] = 6

data[4] = 8

data[5] = 10

data[6] = 12

data[7] = 14

data[8] = 16

data[9] = 18

data[10] = 20

data[11] = 22

data[12] = 24

data[13] = 26

data[14] = 28

data[15] = 30

data[16] = 32

data[17] = 34

data[18] = 36

data[19] = 38

data[20] = 40

data[21] = 42

data[22] = 44

data[23] = 46

data[24] = 48

data[25] = 50

data[26] = 52

data[27] = 54

data[28] = 56

data[29] = 58

data[30] = 60

data[31] = 62

data[32] = 64

data[33] = 66

data[34] = 68

data[35] = 70

data[36] = 72

data[37] = 74

data[38] = 76

data[39] = 78

data[40] = 80

data[41] = 82

data[42] = 84

data[43] = 86

data[44] = 88

data[45] = 90

data[46] = 92

```
data[47] = 94
data[48] = 96
data[49] = 98
data[50] = 100
data[51] = 102
data[52] = 104
data[53] = 106
data[54] = 108
data[55] = 110
data[56] = 112
data[57] = 114
data[58] = 116
data[59] = 118
data[60] = 120
data[61] = 122
data[62] = 124
data[63] = 126
data[64] = 128
data[65] = 130
data[66] = 132
data[67] = 134
data[68] = 136
data[69] = 138
data[70] = 140
data[71] = 142
data[72] = 144
data[73] = 146
data[74] = 148
data[75] = 150
data[76] = 152
data[77] = 154
data[78] = 156
data[79] = 158
data[80] = 160
data[81] = 162
data[82] = 164
data[83] = 166
data[84] = 168
data[85] = 170
data[86] = 172
data[87] = 174
data[88] = 176
data[89] = 178
data[90] = 180
data[91] = 182
data[92] = 184
data[93] = 186
data[94] = 188
data[95] = 190
```

data[96] = 192
data[97] = 194
data[98] = 196
data[99] = 198
Low_bin_search
search 80 index 40
Low_bin_search
search 78 index 39
Low_bin_search
search 76 index 38
Low_bin_search
search 74 index 37
Low_bin_search
search 72 index 36
Low_bin_search
search 70 index 35
Low_bin_search
search 68 index 34
Low_bin_search
search 66 index 33
Low_bin_search
search 64 index 32
Low_bin_search
search 62 index 31
Low_bin_search
search 60 index 30
Low_bin_search
search 58 index 29
Low_bin_search
search 56 index 28
Low_bin_search
search 54 index 27
Low_bin_search
search 52 index 26
Low_bin_search
search 50 index 25
Low_bin_search
search 48 index 24
Low_bin_search
search 46 index 23
Low_bin_search
search 44 index 22
Low_bin_search
search 42 index 21
Low_bin_search
search 40 index 20
Low_bin_search
search 38 index 19
Low_bin_search

```
search 36 index 18
Low_bin_search
search 34 index 17
Low_bin_search
search 32 index 16
Low_bin_search
search 30 index 15
Low_bin_search
search 28 index 14
Low_bin_search
search 26 index 13
Low_bin_search
search 24 index 12
Low_bin_search
search 22 index 11
Low_bin_search
search 20 index 10
Low_bin_search
search 18 index 9
Low_bin_search
search 16 index 8
Low_bin_search
search 14 index 7
Low_bin_search
search 12 index 6
Low_bin_search
search 10 index 5
Low_bin_search
search 8 index 4
Low_bin_search
search 6 index 3
Low_bin_search
search 4 index 2
Low_bin_search
search 2 index 1
Test low_bin_nb_arithmetic
data[0] = 0
data[1] = 2
data[2] = 4
data[3] = 6
data[4] = 8
data[5] = 10
data[6] = 12
data[7] = 14
data[8] = 16
data[9] = 18
data[10] = 20
data[11] = 22
data[12] = 24
```



```
data[13] = 26
data[14] = 28
data[15] = 30
data[16] = 32
data[17] = 34
data[18] = 36
data[19] = 38
data[20] = 40
data[21] = 42
data[22] = 44
data[23] = 46
data[24] = 48
data[25] = 50
data[26] = 52
data[27] = 54
data[28] = 56
data[29] = 58
data[30] = 60
data[31] = 62
data[32] = 64
data[33] = 66
data[34] = 68
data[35] = 70
data[36] = 72
data[37] = 74
data[38] = 76
data[39] = 78
data[40] = 80
data[41] = 82
data[42] = 84
data[43] = 86
data[44] = 88
data[45] = 90
data[46] = 92
data[47] = 94
data[48] = 96
data[49] = 98
data[50] = 100
data[51] = 102
data[52] = 104
data[53] = 106
data[54] = 108
data[55] = 110
data[56] = 112
data[57] = 114
data[58] = 116
data[59] = 118
data[60] = 120
data[61] = 122
```

```
data[62] = 124
data[63] = 126
data[64] = 128
data[65] = 130
data[66] = 132
data[67] = 134
data[68] = 136
data[69] = 138
data[70] = 140
data[71] = 142
data[72] = 144
data[73] = 146
data[74] = 148
data[75] = 150
data[76] = 152
data[77] = 154
data[78] = 156
data[79] = 158
data[80] = 160
data[81] = 162
data[82] = 164
data[83] = 166
data[84] = 168
data[85] = 170
data[86] = 172
data[87] = 174
data[88] = 176
data[89] = 178
data[90] = 180
data[91] = 182
data[92] = 184
data[93] = 186
data[94] = 188
data[95] = 190
data[96] = 192
data[97] = 194
data[98] = 196
data[99] = 198
search 80 index 40
search 78 index 39
search 76 index 38
search 74 index 37
search 72 index 36
search 70 index 35
search 68 index 34
search 66 index 33
search 64 index 32
search 62 index 31
search 60 index 30
```

search 58 index 29
search 56 index 28
search 54 index 27
search 52 index 26
search 50 index 25
search 48 index 24
search 46 index 23
search 44 index 22
search 42 index 21
search 40 index 20
search 38 index 19
search 36 index 18
search 34 index 17
search 32 index 16
search 30 index 15
search 28 index 14
search 26 index 13
search 24 index 12
search 22 index 11
search 20 index 10
search 18 index 9
search 16 index 8
search 14 index 7
search 12 index 6
search 10 index 5
search 8 index 4
search 6 index 3
search 4 index 2
search 2 index 1

Test low_bin_mask

data[0] = 0
data[1] = 2
data[2] = 4
data[3] = 6
data[4] = 8
data[5] = 10
data[6] = 12
data[7] = 14
data[8] = 16
data[9] = 18
data[10] = 20
data[11] = 22
data[12] = 24
data[13] = 26
data[14] = 28
data[15] = 30
data[16] = 32
data[17] = 34
data[18] = 36

```
data[19] = 38
data[20] = 40
data[21] = 42
data[22] = 44
data[23] = 46
data[24] = 48
data[25] = 50
data[26] = 52
data[27] = 54
data[28] = 56
data[29] = 58
data[30] = 60
data[31] = 62
data[32] = 64
data[33] = 66
data[34] = 68
data[35] = 70
data[36] = 72
data[37] = 74
data[38] = 76
data[39] = 78
data[40] = 80
data[41] = 82
data[42] = 84
data[43] = 86
data[44] = 88
data[45] = 90
data[46] = 92
data[47] = 94
data[48] = 96
data[49] = 98
data[50] = 100
data[51] = 102
data[52] = 104
data[53] = 106
data[54] = 108
data[55] = 110
data[56] = 112
data[57] = 114
data[58] = 116
data[59] = 118
data[60] = 120
data[61] = 122
data[62] = 124
data[63] = 126
data[64] = 128
data[65] = 130
data[66] = 132
data[67] = 134
```

```
data[68] = 136
data[69] = 138
data[70] = 140
data[71] = 142
data[72] = 144
data[73] = 146
data[74] = 148
data[75] = 150
data[76] = 152
data[77] = 154
data[78] = 156
data[79] = 158
data[80] = 160
data[81] = 162
data[82] = 164
data[83] = 166
data[84] = 168
data[85] = 170
data[86] = 172
data[87] = 174
data[88] = 176
data[89] = 178
data[90] = 180
data[91] = 182
data[92] = 184
data[93] = 186
data[94] = 188
data[95] = 190
data[96] = 192
data[97] = 194
data[98] = 196
data[99] = 198
search 80 index 40
search 78 index 39
search 76 index 38
search 74 index 37
search 72 index 36
search 70 index 35
search 68 index 34
search 66 index 33
search 64 index 32
search 62 index 31
search 60 index 30
search 58 index 29
search 56 index 28
search 54 index 27
search 52 index 26
search 50 index 25
search 48 index 24
```

search 46 index 23
search 44 index 22
search 42 index 21
search 40 index 20
search 38 index 19
search 36 index 18
search 34 index 17
search 32 index 16
search 30 index 15
search 28 index 14
search 26 index 13
search 24 index 12
search 22 index 11
search 20 index 10
search 18 index 9
search 16 index 8
search 14 index 7
search 12 index 6
search 10 index 5
search 8 index 4
search 6 index 3
search 4 index 2
search 2 index 1

Test low_bin_nb_4x

data[0] = 0
data[1] = 2
data[2] = 4
data[3] = 6
data[4] = 8
data[5] = 10
data[6] = 12
data[7] = 14
data[8] = 16
data[9] = 18
data[10] = 20
data[11] = 22
data[12] = 24
data[13] = 26
data[14] = 28
data[15] = 30
data[16] = 32
data[17] = 34
data[18] = 36
data[19] = 38
data[20] = 40
data[21] = 42
data[22] = 44
data[23] = 46
data[24] = 48

```
data[25] = 50
data[26] = 52
data[27] = 54
data[28] = 56
data[29] = 58
data[30] = 60
data[31] = 62
data[32] = 64
data[33] = 66
data[34] = 68
data[35] = 70
data[36] = 72
data[37] = 74
data[38] = 76
data[39] = 78
data[40] = 80
data[41] = 82
data[42] = 84
data[43] = 86
data[44] = 88
data[45] = 90
data[46] = 92
data[47] = 94
data[48] = 96
data[49] = 98
data[50] = 100
data[51] = 102
data[52] = 104
data[53] = 106
data[54] = 108
data[55] = 110
data[56] = 112
data[57] = 114
data[58] = 116
data[59] = 118
data[60] = 120
data[61] = 122
data[62] = 124
data[63] = 126
data[64] = 128
data[65] = 130
data[66] = 132
data[67] = 134
data[68] = 136
data[69] = 138
data[70] = 140
data[71] = 142
data[72] = 144
data[73] = 146
```

```
data[74] = 148
data[75] = 150
data[76] = 152
data[77] = 154
data[78] = 156
data[79] = 158
data[80] = 160
data[81] = 162
data[82] = 164
data[83] = 166
data[84] = 168
data[85] = 170
data[86] = 172
data[87] = 174
data[88] = 176
data[89] = 178
data[90] = 180
data[91] = 182
data[92] = 184
data[93] = 186
data[94] = 188
data[95] = 190
data[96] = 192
data[97] = 194
data[98] = 196
data[99] = 198
search 80 index 40
search 78 index 39
search 76 index 38
search 74 index 37
search 72 index 36
search 70 index 35
search 68 index 34
search 66 index 33
search 64 index 32
search 62 index 31
search 60 index 30
search 58 index 29
search 56 index 28
search 54 index 27
search 52 index 26
search 50 index 25
search 48 index 24
search 46 index 23
search 44 index 22
search 42 index 21
search 40 index 20
search 38 index 19
search 36 index 18
```


search 34 index 17
search 32 index 16
search 30 index 15
search 28 index 14
search 26 index 13
search 24 index 12
search 22 index 11
search 20 index 10
search 18 index 9
search 16 index 8
search 14 index 7
search 12 index 6
search 10 index 5
search 8 index 4
search 6 index 3
search 4 index 2
search 2 index 1

db5242test.c: In function 'main':

db5242test.c:663:43: error: expected ';' before 'int64_t'

```
663 |     printf("Test low_bin_nb_simd simd \n")
```

```
    |                               ^  
    |                               ;
```

```
664 |     int64_t arraysize;
```

```
    |     ~~~~~
```

Test low_bin_nb_4x

data[0] = 0

data[1] = 2

data[2] = 4

data[3] = 6

data[4] = 8

data[5] = 10

data[6] = 12

data[7] = 14

data[8] = 16

data[9] = 18

data[10] = 20

data[11] = 22

data[12] = 24

data[13] = 26

data[14] = 28

data[15] = 30

data[16] = 32

data[17] = 34

data[18] = 36

data[19] = 38

data[20] = 40

data[21] = 42

data[22] = 44

data[23] = 46

```
data[24] = 48
data[25] = 50
data[26] = 52
data[27] = 54
data[28] = 56
data[29] = 58
data[30] = 60
data[31] = 62
data[32] = 64
data[33] = 66
data[34] = 68
data[35] = 70
data[36] = 72
data[37] = 74
data[38] = 76
data[39] = 78
data[40] = 80
data[41] = 82
data[42] = 84
data[43] = 86
data[44] = 88
data[45] = 90
data[46] = 92
data[47] = 94
data[48] = 96
data[49] = 98
data[50] = 100
data[51] = 102
data[52] = 104
data[53] = 106
data[54] = 108
data[55] = 110
data[56] = 112
data[57] = 114
data[58] = 116
data[59] = 118
data[60] = 120
data[61] = 122
data[62] = 124
data[63] = 126
data[64] = 128
data[65] = 130
data[66] = 132
data[67] = 134
data[68] = 136
data[69] = 138
data[70] = 140
data[71] = 142
data[72] = 144
```

```
data[73] = 146
data[74] = 148
data[75] = 150
data[76] = 152
data[77] = 154
data[78] = 156
data[79] = 158
data[80] = 160
data[81] = 162
data[82] = 164
data[83] = 166
data[84] = 168
data[85] = 170
data[86] = 172
data[87] = 174
data[88] = 176
data[89] = 178
data[90] = 180
data[91] = 182
data[92] = 184
data[93] = 186
data[94] = 188
data[95] = 190
data[96] = 192
data[97] = 194
data[98] = 196
data[99] = 198
search 80 index 40
search 78 index 39
search 76 index 38
search 74 index 37
search 72 index 36
search 70 index 35
search 68 index 34
search 66 index 33
search 64 index 32
search 62 index 31
search 60 index 30
search 58 index 29
search 56 index 28
search 54 index 27
search 52 index 26
search 50 index 25
search 48 index 24
search 46 index 23
search 44 index 22
search 42 index 21
search 40 index 20
search 38 index 19
```

search 36 index 18
search 34 index 17
search 32 index 16
search 30 index 15
search 28 index 14
search 26 index 13
search 24 index 12
search 22 index 11
search 20 index 10
search 18 index 9
search 16 index 8
search 14 index 7
search 12 index 6
search 10 index 5
search 8 index 4
search 6 index 3
search 4 index 2
search 2 index 1

Test the band_join (0/2)

Test Failed: 28416 != 0

Test the band_join with SIMD (0/2)

Test Failed: 28416 != 0

Test the low bin search and it should always pass (2/2)

Compilation with mavn2 (2/2)

Compilation with mavn512f (2/2)

Test the low_bin_nb_arithmetic (2/2)

Test the low_bin_nb_mask (2/2)

Test the low_bin_nb_4x (2/2)

Test the low_bin_nb_simd (2/2)


```
1  /*
2   CSE 5242 Project 2, Fall 2023
3
4   See class project handout for more extensive documentation.
5
6   https://stackoverflow.com/questions/19068705/undefined-reference-when-calling-inline-function
7  */
8
9  #include <stdlib.h>
10 #include <stdio.h>
11 #include <stdint.h>
12 #include <unistd.h>
13 #include <string.h>
14 #include <sys/ioctl.h>
15 #include <sys/time.h>
16 #include <asm/unistd.h>
17 #include <immintrin.h>
18
19 /* uncomment out the following DEBUG line for debug info, for experiment comment the DEBUG line
20  */
21 // #define DEBUG
22
23 /* compare two int64_t values - for use with qsort */
24 static int compare(const void *p1, const void *p2)
25 {
26     int a,b;
27     a = *(int64_t *)p1;
28     b = *(int64_t *)p2;
29     if (a<b) return -1;
30     if (a==b) return 0;
31     return 1;
32 }
33
34 /* initialize searches and data - data is sorted and searches is a random permutation of data */
35 int init(int64_t* data, int64_t* searches, int count)
36 {
37     for(int64_t i=0; i<count; i++){
38         searches[i] = random();
39         data[i] = searches[i]+1;
40     }
41     qsort(data,count,sizeof(int64_t),compare);
42 }
43
44 /* initialize outer probes of band join */
45 int band_init(int64_t* outer, int64_t size)
```

```

46 {
47     for(int64_t i=0; i<size; i++){
48         outer[i] = random();
49     }
50 }
51
52 inline int64_t simple_binary_search(int64_t* data, int64_t size, int64_t target)
53 {
54     int64_t left=0;
55     int64_t right=size;
56     int64_t mid;
57
58     while(left<=right) {
59         mid = (left + right)/2; /* ignore possibility of overflow of left+right */
60         if (data[mid]==target) return mid;
61         if (data[mid]<target) left=mid+1;
62         else right = mid-1;
63     }
64     return -1; /* no match */
65 }
66
67 inline int64_t low_bin_search(int64_t* data, int64_t size, int64_t target)
68 {
69     /* this binary search variant
70      (a) does only one comparison in the inner loop
71      (b) doesn't require an exact match; instead it returns the index of the first key >= the search key.
72          That's good in a DB context where we might be doing a range search, and using binary search to
73          identify the first key in the range.
74      (c) If the search key is bigger than all keys, it returns size.
75     */
76     int64_t left=0;
77     int64_t right=size;
78     int64_t mid;
79
80     printf("Low_bin_search\n");
81     while(left<right) {
82         mid = (left + right)/2; /* ignore possibility of overflow of left+right */
83         if (data[mid]>=target)
84             right=mid;
85         else
86             left=mid+1;
87     }
88     return right;
89 }
90
91 //define ARDEBUB
92 inline int64_t low_bin_nb_arithmetic(int64_t* data, int64_t size, int64_t target)
93 {
94     /* this binary search variant

```

```

95     (a) does no comparisons in the inner loop by using multiplication and addition to convert control
dependencies
96     to data dependencies
97     (b) doesn't require an exact match; instead it returns the index of the first key >= the search key.
98     That's good in a DB context where we might be doing a range search, and using binary search to
99     identify the first key in the range.
100    (c) If the search key is bigger than all keys, it returns size.
101    */
102    int64_t left=0;
103    int64_t right=size;
104    int64_t mid;
105
106    #ifdef ARDEBUG
107        printf("low_bin_nb_arith\n");
108    #endif
109    //0^1 = 1 and 1^1 = 0
110    while(left<right) {
111        mid = (left + right) / 2; //get middle
112        int64_t flag = data[mid] >= target; //data[mid] >= target ? 1 : 0 (if target value is left side of the
middle index flag = 1 else 0)
113        left = flag * left + (flag^1) * (mid+1); // if flag == 1 assign previous left value (left index stays same)
else assign mid+1 value to divide the array into half
114        right = flag * mid + (flag^1) * right; // if flag == 1 assign mid value to divide array into half else assign
previous right value (right index stays same)
115        #ifdef ARDEBUG
116            printf("mid: %d left: %d right: %d\n", mid, left, right);
117        #endif
118    }
119    return right;
120 }
121
122 //define MASTDEBUG
123 inline int64_t low_bin_nb_mask(int64_t* data, int64_t size, int64_t target)
124 {
125     /* this binary search variant
126     (a) does no comparisons in the inner loop by using bit masking operations to convert control
dependencies
127     to data dependencies
128     (b) doesn't require an exact match; instead it returns the index of the first key >= the search key.
129     That's good in a DB context where we might be doing a range search, and using binary search to
130     identify the first key in the range.
131     (c) If the search key is bigger than all keys, it returns size.
132     */
133     int64_t left=0;
134     int64_t right=size;
135     int64_t mid;
136
137     #ifdef MASTDEBUG
138         printf("low_bin_nb_mask\n");

```



```

139 #endif
140 // ~(-1) = 0 and ~(0) = -1
141 //-1 & 3 = 3 and 0 & 3 = 0
142 //0 | 3 = 3
143 while(left<right) {
144     //if-else using bitwise source: https://stackoverflow.com/questions/3798601/conditional-using-bitwise-operator
145     mid = (left + right) / 2; //get middle
146     int64_t flag = (data[mid] >= target)-1; //data[mid] >= target ? 0 : -1 (if target value is left side of
middle index flag = 0 else =1)
147     left = (flag & (mid+1)) | (~flag & left); // If flag == -1 assign mid+1 to divide the array into half else
assign previous left value (left index stays same)
148     right = (flag & right) | (~flag & mid); // if flag == -1 assign previous right value (right index stay same)
else assing mid to divide the array into half
149     #ifdef MASTDEBUG
150         printf("mid: %d left: %d right: %d\n", mid, left, right);
151     #endif
152     //int64_t flag = (data[mid] >= target); //data[mid] >= target ? 1 : 0
153     // left = ((flag-1) & (mid+1)) | ((flag^1)-1 & left);
154     // right = ((flag-1) & right) | ((flag^1)-1 & mid);
155 }
156 return right;
157 }
158
159 //define XDEBUG
160 inline void low_bin_nb_4x(int64_t* data, int64_t size, int64_t* targets, int64_t* right)
161 {
162     /* this binary search variant
163     (a) does no comparisons in the inner loop by using bit masking operations instead
164     (b) doesn't require an exact match; instead it returns the index of the first key >= the search key.
165         That's good in a DB context where we might be doing a range search, and using binary search to
166         identify the first key in the range.
167     (c) If the search key is bigger than all keys, it returns size.
168     (d) does 4 searches at the same time in an interleaved fashion, so that an out-of-order processor
169         can make progress even when some instructions are still waiting for their operands to be ready.
170
171     Note that we're using the result array as the "right" elements in the search so no need for a return
statement
172     */
173     int64_t left[4]={0,0,0,0};
174     int64_t mid[4];
175     right[0]=right[1]=right[2]=right[3]=size;
176
177     #ifdef XDEBUG
178         printf("low_bin_nb_4x\n");
179     #endif
180     int64_t bit_size = 4;
181     int cnt = 0;
182     //outer loop to check termination condition like above (e.g left < right)

```

```

183 while (left[0]<right[0] || left[1]<right[1] || left[2]<right[2] || left[3]<right[3]) {
184     int i;
185     #ifdef XDEBUG
186         printf("cnt: %d\n", cnt);
187     #endif
188     //inner loop to perform 4 search concurrently
189     for (i = 0; i < bit_size; i++) {
190         //if-else using bitwise source: https://stackoverflow.com/questions/3798601/conditional-using-
        bitwise-operator
191         mid[i] = (left[i] + right[i]) / 2; //get middle
192         int64_t flag = (data[mid[i]] >= targets[i]) - 1; //data[mid] >= target ? 0 : -1 (if target value is left side
        of middle index flag = 0 else =1)
193         left[i] = (flag & (mid[i]+1)) | (~flag & left[i]); // If flag == -1 assign mid+1 to divide the array into half
        else assign previous left value (left index stays same)
194         right[i] = (flag & right[i]) | (~flag & mid[i]); // if flag == -1 assign previous right value (right index
        stay same) else assing mid to divide the array into half
195         #ifdef XDEBUG
196             printf("mid: %d left: %d right: %d\n", mid[i], left[i], right[i]);
197         #endif
198
199         //int64_t flag = (data[mid[i]] > targets[i]); //data[mid] >= target ? 1 : 0
200         // left[i] = ((flag-1) & (mid[i]+1)) | ((flag^1)-1 & left[i]);
201         // right[i] = ((flag-1) & right[i]) | ((flag^1)-1 & mid[i]);
202         //break;
203     }
204     cnt+=1;
205 }
206 }
207
208
209 /* The following union type is handy to output the contents of AVX512 data types */
210 union int8x4 {
211     __m256i a;
212     int64_t b[4];
213 };
214
215 void printavx(char* name, __m256i v) {
216     union int8x4 n;
217
218     n.a=v;
219     printf("Value in %s is [%ld %ld %ld %ld ]\n",name,n.b[0],n.b[1],n.b[2],n.b[3]);
220 }
221
222 /*
223  * Optinal for using AVX-512
224
225 union int8x8 {
226     __m512i a;
227     int64_t b[8];

```

```

228 };
229
230 void printavx512(char* name, __m512i v) {
231     union int8x4 n;
232
233     n.a=v;
234     printf("Value in %s is [%ld %ld %ld %ld %ld %ld %ld %ld %ld %ld %ld %ld %ld %ld %ld %ld]\n",name,n.b[0],n.b[1],n.b[2],n.b[3]);
235 }
236
237 */
238
239 //define SIMDDEBUG
240 inline void low_bin_nb_simd(int64_t* data, int64_t size, __m256i target, __m256i* result)
241 {
242     /* this binary search variant
243     (a) does no comparisons in the inner loop by using bit masking operations instead
244     (b) doesn't require an exact match; instead it returns the index of the first key >= the search key.
245         That's good in a DB context where we might be doing a range search, and using binary search to
246         identify the first key in the range.
247     (c) If the search key is bigger than all keys, it returns size.
248     (d) does 4 searches at the same time using AVX2 intrinsics
249
250     See https://www.intel.com/content/www/us/en/docs/cpp-compiler/developer-guide-reference/2021-8/intrinsics-for-avx2.html
251     for documentation of the AVX512 intrinsics
252
253     Note that we're using the result array as the "right" elements in the search, and that searchkey is
254     being passed
255     as an __m256i value rather than via a pointer.
256
257     */
258     __m256i aleft = _mm256_set1_epi64x(0);
259     __m256i aright = _mm256_set1_epi64x(size);
260     __m256i amid;
261
262     __m256i ones = _mm256_set1_epi64x(1);
263     __m256i amask;
264     __m256i datavec;
265
266     #ifdef SIMDDEBUG
267     printf("low_bin_nb_simd\n");
268     #endif
269
270     //AVX Intrinsic Guide: https://www.intel.com/content/www/us/en/docs/intrinsics-guide/index.html
271     __m256i cmp = _mm256_cmpgt_epi64(aright, aleft);
272     //Check if all elements of vector is 0 using testz: https://stackoverflow.com/questions/32072169/could-i-compare-to-zero-register-in-avx-correctly + https://stackoverflow.com/a/32120039
273     //Check all elements in cmp are 0 (if value of cmp is 0 it means left > right else -1)
274     while(_mm256_testz_si256(cmp, cmp) == 0){
275         amid = _mm256_srli_epi64(_mm256_add_epi64(aleft, aright), 1); //get middle

```

```

274   _mm256i amid_one = _mm256_add_epi64(amid, ones); //get middle+1 vector
275   //Extract matching values of data using AVX index vector using i64gather:
https://stackoverflow.com/questions/51128005/what-do-you-do-without-fast-gather-and-scatter-in-
avx2-instructions
276   datavec = _mm256_i64gather_epi64((long long*)data, amid, sizeof(int64_t));
277
278   // inlining failed error for >= operation
279   //__mmask8 flag = _mm256_cmpgqe_epi64_mask(datavec, target);
280   //Combine > operation and = operation to have >= operation: https://www.splashlearn.com/math-
vocabulary/greater-than-or-equal-to
281   _mm256i cmpgt = _mm256_cmpgt_epi64(datavec, target);
282   _mm256i cmpeq = _mm256_cmpeq_epi64(datavec, target);
283   _mm256i flag = _mm256_or_si256(cmpgt, cmpeq);
284
285   //if-else using bitwise source: https://stackoverflow.com/questions/3798601/conditional-using-
bitwise-operator
286   aleft = _mm256_or_si256(_mm256_and_si256(flag, aleft), _mm256_andnot_si256(flag, amid_one)); // if
falg == -1 assign left value (left index stays same) else assing mid+1 to divide the array into half.
287   aright = _mm256_or_si256(_mm256_and_si256(flag, amid), _mm256_andnot_si256(flag, aright)); // if
flag == -1 assign mid value to divide array into half else assign previous right value (rihgt index stays
same)
288
289   #ifdef SIMDDEBUG
290       printavx("amid", amid);
291       printavx("flag", flag);
292       printavx("aright", aright);
293       printavx("aleft", aleft);
294       printf("\n\n");
295   #endif
296   //break;
297   cmp = _mm256_cmpgt_epi64(aright, aleft);
298   }
299   //aright = _mm256_sub_epi64(aright, ones);
300   *result = aright;
301   }
302
303   void bulk_bin_search(int64_t* data, int64_t size, int64_t* searchkeys, int64_t numsearches, int64_t*
results, int repeats)
304   {
305       for(int j=0; j<repeats; j++) {
306           /* Function to test a large number of binary searches
307
308               we might need repeats>1 to make sure the events we're measuring are not dominated by various
309               overheads, particularly for small values of size and/or numsearches
310
311               we assume that we want exactly "size" searches, where "size" is the length if the searchkeys array
312               */
313           for(int64_t i=0; i<numsearches; i++) {
314   #ifdef DEBUG

```

```

315     printf("Searching for %ld...\n",searchkeys[i]);
316 #endif
317
318     // Uncomment one of the following to measure it
319     //results[i] = low_bin_search(data,size,searchkeys[i]);
320     //results[i] = low_bin_nb_arithmetic(data,size,searchkeys[i]);
321     results[i] = low_bin_nb_mask(data,size,searchkeys[i]);
322
323 #ifdef DEBUG
324     printf("Result is %ld\n",results[i]);
325 #endif
326 }
327 }
328 }
329
330 void bulk_bin_search_4x(int64_t* data, int64_t size, int64_t* searchkeys, int64_t numsearches, int64_t*
results, int repeats)
331 {
332     register __m256i searchkey_4x;
333
334     for(int j=0; j<repeats; j++) {
335         /* Function to test a large number of binary searches using one of the 8x routines
336
337         we might need repeats>1 to make sure the events we're measuring are not dominated by various
338         overheads, particularly for small values of size and/or numsearches
339
340         we assume that we want exactly "size" searches, where "size" is the length if the searchkeys array
341         */
342         int64_t extras = numsearches % 4;
343         for(int64_t i=0;i<numsearches-extras; i+=4) {
344             #ifdef DEBUG
345                 printf("Searching for %ld %ld %ld %ld ...\\n",
346                     searchkeys[i],searchkeys[i+1],searchkeys[i+2],searchkeys[i+3]);
347             #endif
348
349             // Uncomment one of the following depending on which routine you want to profile
350
351             // Algorithm A
352             //low_bin_nb_4x(data,size,&searchkeys[i],&results[i]);
353
354             // Algorithm B
355             searchkey_4x = _mm256_loadu_si256((__m256i *)&searchkeys[i]);
356             low_bin_nb_simd(data,size,searchkey_4x,(__m256i *)&results[i]);
357
358             #ifdef DEBUG
359                 printf("Result is %ld %ld %ld %ld ...\\n",
360                     results[i],results[i+1],results[i+2],results[i+3]);
361             #endif
362         }

```

```

363  /* a little bit more work if numsearches is not a multiple of 8 */
364  for(int64_t i=numsearches-extras;i<numsearches; i++) {
365
366      results[i] = low_bin_nb_mask(data,size,searchkeys[i]);
367
368  }
369
370  }
371  }
372
373
374  int64_t band_join(int64_t* inner, int64_t inner_size, int64_t* outer, int64_t outer_size, int64_t*
inner_results, int64_t* outer_results, int64_t result_size, int64_t bound)
375  {
376      /* In a band join we want matches within a range of values. If p is the probe value from the outer
table, then all
377      records in the inner table with a key in the range [p-bound,p+bound] inclusive should be part of
the result.
378
379      Results are returned via two arrays. outer_results stores the index of the outer table row that
matches, and
380      inner_results stores the index of the inner table row that matches. result_size tells you the size of
the
381      output array that has been allocated. You should make sure that you don't exceed this size. If there
are
382      more results than can fit in the result arrays, then return early with just a prefix of the results in the
result
383      arrays. The return value of the function should be the number of output results.
384
385      */
386      //Declaring arrays to store left and right index of the range.
387      int64_t *leftIndexArray = malloc(outer_size * sizeof(int64_t));
388      int64_t *rightIndexArray = malloc(outer_size * sizeof(int64_t));
389      if (!leftIndexArray || !rightIndexArray) {
390          free(leftIndexArray);
391          free(rightIndexArray);
392          // Memory allocation failed, so we return -1
393          return -1;
394      }
395      //Calculating below values to find the multiplier and remainder.
396      int multiplier = outer_size/4;
397      int remainder = outer_size%4;
398
399      int temp_1=0;
400
401      while (temp_1<multiplier){
402          // printf("Inside multiplier with i->%d\n",temp_1);
403          //using the below arrays to store the lower bound values temporarily after using low_bin_nb_mask
function.

```

```

404     int64_t lower_limit[4];
405     int64_t upper_limit[4];
406
407     for (int i = 0; i < 4; i++) {
408         lower_limit[i] = outer[4 * temp_1 + i] - bound;
409         upper_limit[i] = outer[4 * temp_1 + i] + bound;
410     }
411     int64_t left_temp_Array[4];
412     int64_t right_temp_Array[4];
413
414     low_bin_nb_4x(inner, inner_size, lower_limit, left_temp_Array);
415     low_bin_nb_4x(inner, inner_size, upper_limit, right_temp_Array);
416
417     for (int i = 0; i < 4; i++) {
418         leftIndexArray[4 * temp_1 + i] = left_temp_Array[i];
419         rightIndexArray[4 * temp_1 + i] = right_temp_Array[i];
420     }
421     // printf("Ending multiplier with i->%d\n", temp_1);
422     temp_1 += 1;
423
424 }
425
426 for (int i = 0; i < remainder; i++) {
427     // printf("Entering remainder with i->%d\n", i);
428     // printf("lower-> %d\n", outer[i] - bound);
429     // printf("upper-> %d\n", outer[i] + bound);
430     leftIndexArray[4 * temp_1 + i] = low_bin_nb_mask(inner, inner_size, outer[4 * temp_1 + i] - bound); //array
containing left index in the range.
431     rightIndexArray[4 * temp_1 + i] = low_bin_nb_mask(inner, inner_size, outer[4 * temp_1 + i] + bound); //array
containing right index of the range.
432     // printf("leftindex->%d \t\n rightindex->%d\t\n", leftIndexArray[i], rightIndexArray[i]);
433     // printf("Exiting reminder loop with i->%d\n", i);
434 } printf("\n");
435
436
437
438 /* ONLY WITH low_bin_nb_mask
439 for (int i = 0; i < outer_size; i++) {
440     // printf("lower-> %d\n", outer[i] - bound);
441     // printf("upper-> %d\n", outer[i] + bound);
442     leftIndexArray[i] = low_bin_nb_mask(inner, inner_size, outer[i] - bound); //array containing left index in
the range.
443     rightIndexArray[i] = low_bin_nb_mask(inner, inner_size, outer[i] + bound); //array containing right index
of the range.
444     // printf("leftindex->%d \t\n rightindex->%d\t\n", leftIndexArray[i], rightIndexArray[i]);
445 } printf("\n");
446 */
447 // for (int i = 0; i < outer_size; i++) {
448 // printf("leftIndexArray %d->%d and rightIndexArray %d -

```

```

>%d\n",i,leftIndexArray[i],i,rightIndexArray[i]);
449 // }
450
451 int temp=0;
452
453 for(int i = 0; i<outer_size;i++){
454     // printf("Entering inside the merging part with i->%d\n",i);
455     int left_index = leftIndexArray[i];
456     int right_index = rightIndexArray[i];
457     //if left index and right index are same then we skip the iteration because there are no elements.
458     if (left_index==right_index){continue;}
459
460     while ((temp<result_size) && (left_index<right_index)){
461         // printf("Inside while for outer element index %d\n",i);
462         outer_results[temp]=i;
463         inner_results[temp]=left_index;
464         // printf("adding values (%d,%d) in outer and inner result\n",outer_results[i],inner_results[i]);
465         left_index+=1;
466         temp+=1;
467     }
468     //break if we reach the result size.
469     if (temp==result_size){break;}
470     // printf("Exiting merging part with i->%d\n",i);
471
472 }
473
474
475 // for (int i = 0; i<temp;i++){
476 //     printf("(OuterResult,InnerResult)->(%ld,%ld)\n",outer_results[i],inner_results[i]);
477 // }
478
479 free(leftIndexArray);
480 free(rightIndexArray);
481
482 return temp;
483
484 }
485
486 int64_t band_join_simd(int64_t* inner, int64_t inner_size, int64_t* outer, int64_t outer_size, int64_t*
inner_results, int64_t* outer_results, int64_t result_size, int64_t bound)
487 {
488     /* In a band join we want matches within a range of values. If p is the probe value from the outer
table, then all
489     records in the inner table with a key in the range [p-bound,p+bound] inclusive should be part of
the result.
490
491     Results are returned via two arrays. outer_results stores the index of the outer table row that
matches, and
492     inner_results stores the index of the inner table row that matches. result_size tells you the size of

```


the

```
493     output array that has been allocated. You should make sure that you don't exceed this size. If there
are
494     more results than can fit in the result arrays, then return early with just a prefix of the results in the
result
495     arrays. The return value of the function should be the number of output results.
496
497     To do the binary search, you could use the low_bin_nb_simd you just implemented to search for the
lower bounds in parallel
498
499     Once you've found the lower bounds, do the following for each of the 4 search keys in turn:
500     scan along the sorted inner array, generating outputs for each match, and making sure not to
exceed the output array bounds.
501
502     This inner scanning code does not have to use SIMD.
503     */
504
505
506
507 //Using similiar approach from the above function.
508
509 // Create and allocate space to store left and right result indexes
510 int64_t *leftIndexArray = malloc(outer_size * sizeof(int64_t));
511 int64_t *rightIndexArray = malloc(outer_size * sizeof(int64_t));
512 if (!leftIndexArray || !rightIndexArray) {
513     free(leftIndexArray);
514     free(rightIndexArray);
515     return -1;
516 }
517
518 // Calculate multiplier and remainder to find number of simd function instances and individual mask
function call instances
519 int multiplier = outer_size/4;
520 int remainder = outer_size%4;
521
522 int temp_1=0;
523 while (temp_1<multiplier){
524     // compute lower and upper bound vectors using AVX functions
525     __m256i outer_vec = _mm256_loadu_si256((__m256i*)&outer[4*temp_1]);
526     __m256i lower_bound_vec = _mm256_sub_epi64(outer_vec, _mm256_set1_epi64x(bound));
527     __m256i upper_bound_vec = _mm256_add_epi64(outer_vec, _mm256_set1_epi64x(bound));
528
529     // Create temporary indexes to store each iteration's results
530     int64_t left_indexes[4];
531     int64_t right_indexes[4];
532
533     // Call binary search function using simd
534     low_bin_nb_simd(inner,inner_size,lower_bound_vec,(__m256i*) left_indexes);
535     low_bin_nb_simd(inner,inner_size,upper_bound_vec,(__m256i*) right_indexes);
```

```

536
537 // Store results in temporary arrays
538 for (int i =0;i<4;i++){
539     leftIndexArray[4*temp_1 + i]=left_indexes[i];
540     rightIndexArray[4*temp_1 + i]=right_indexes[i];
541 }
542 temp_1+=1;
543
544 }
545
546 // Store indexes of remaining cases into the result arrays
547 for(int i = 0; i<remainder;i++){
548     leftIndexArray[4*temp_1 + i]=low_bin_nb_mask(inner, inner_size, outer[4*temp_1 + i]-bound);//array
containing left index in the range.
549     rightIndexArray[4*temp_1 + i]=low_bin_nb_mask(inner,inner_size,outer[4*temp_1 + i]+bound);//array
containing right index of the range.
550 }
551
552 // temp will be the final result size
553 int temp=0;
554
555 for(int i = 0; i<outer_size;i++){
556     int left_index = leftIndexArray[i];
557     int right_index = rightIndexArray[i];
558
559     if (left_index==right_index){continue;}
560
561     // Store indexes into the result format
562     while ((temp<result_size) && (left_index<right_index)){
563         outer_results[temp]=i;
564         inner_results[temp]=left_index;
565         left_index+=1;
566         temp+=1;
567     }
568 }
569
570 if (temp==result_size){break;}
571
572 }
573
574 // Clear index array space
575 free(leftIndexArray);
576 free(rightIndexArray);
577
578 // Return result size
579 return temp;
580
581
582 /*

```

```

583 INITIAL CODE WHICH WAS GIVING INCONSISTENT RESULT. PLEASE SKIP TO END OF COMMENT FOR
THE FINAL CODE.
584 int size = 0;
585
586 for(int i=0; (i<outer_size) && (size<result_size); i+=4){
587     __m256i outer_vec = _mm256_loadu_si256((__m256i*)&outer[i]);
588     __m256i lower_bound_vec = _mm256_sub_epi64(outer_vec, _mm256_set1_epi64x(bound));
589     __m256i upper_bound_vec = _mm256_add_epi64(outer_vec, _mm256_set1_epi64x(bound));
590
591     int64_t left_indexes[4];
592     int64_t right_indexes[4];
593
594     int remainder = outer_size - i;
595     if (remainder >= 4) {
596
597         __m256i left_result[1];
598         __m256i right_result[1];
599
600         low_bin_nb_simd(inner,inner_size,lower_bound_vec,left_result);
601         low_bin_nb_simd(inner,inner_size,upper_bound_vec,right_result);
602
603         _mm256_store_si256((__m256i*)left_indexes, left_result[0]);
604         _mm256_store_si256((__m256i*)right_indexes, right_result[0]);
605
606         for(int j = 0; j < 4; j++) {
607             while((size < result_size) && (left_indexes[j] < right_indexes[j])) {
608                 outer_results[size] = i+j;
609                 inner_results[size] = left_indexes[j];
610                 left_indexes[j]++;
611                 size++;
612             }
613         }
614     } else {
615
616         if (remainder >= 1) {
617             int64_t lower_bound = _mm256_extract_epi64(lower_bound_vec, 0);
618             int64_t upper_bound = _mm256_extract_epi64(upper_bound_vec, 0);
619
620             left_indexes[0] = low_bin_nb_mask(inner,inner_size,lower_bound);
621             right_indexes[0] = low_bin_nb_mask(inner,inner_size,upper_bound);
622         }
623
624         if (remainder >= 2) {
625             int64_t lower_bound = _mm256_extract_epi64(lower_bound_vec, 1);
626             int64_t upper_bound = _mm256_extract_epi64(upper_bound_vec, 1);
627
628             left_indexes[1] = low_bin_nb_mask(inner,inner_size,lower_bound);
629             right_indexes[1] = low_bin_nb_mask(inner,inner_size,upper_bound);
630

```

```

631     }
632
633     if (remainder == 3) {
634         int64_t lower_bound = _mm256_extract_epi64(lower_bound_vec, 2);
635         int64_t upper_bound = _mm256_extract_epi64(upper_bound_vec, 2);
636
637         left_indexes[2] = low_bin_nb_mask(inner, inner_size, lower_bound);
638         right_indexes[2] = low_bin_nb_mask(inner, inner_size, upper_bound);
639     }
640
641     for(int j = 0; j < remainder; j++) {
642         while((size < result_size) && (left_indexes[j] < right_indexes[j])) {
643             outer_results[size] = i+j;
644             inner_results[size] = left_indexes[j];
645             left_indexes[j]++;
646             size++;
647         }
648     }
649
650 }
651
652 }
653 return size;
654 */
655
656 }
657
658 int main(int argc, char *argv[])
659 {
660     long long counter;
661     int64_t arraysize, outer_size, result_size;
662     int64_t bound;
663     int64_t *data, *queries, *results;
664     int ret;
665     struct timeval before, after;
666     int repeats;
667     int64_t total_results;
668
669     // band-join arrays
670     int64_t *outer, *outer_results, *inner_results;
671
672
673     if (argc >= 5)
674     {
675         arraysize = atoi(argv[1]);
676         outer_size = atoi(argv[2]);
677         result_size = atoi(argv[3]);
678         bound = atoi(argv[4]);
679     }

```

```

680     else
681     {
682         fprintf(stderr, "Usage: db5242 inner_size outer_size result_size bound <repeats>\n");
683         exit(EXIT_FAILURE);
684     }
685
686     if (argc >= 6)
687         repeats = atoi(argv[5]);
688     else
689     {
690         repeats=1;
691     }
692 // printf("InsideMain and bound->%d\n",bound);
693
694     /* allocate the array and the queries for searching */
695     ret=posix_memalign((void**) &data,64,array_size*sizeof(int64_t));
696     if (ret)
697     {
698         fprintf(stderr, "Memory allocation error.\n");
699         exit(EXIT_FAILURE);
700     }
701     ret=posix_memalign((void**) &queries,64,array_size*sizeof(int64_t));
702     if (ret)
703     {
704         fprintf(stderr, "Memory allocation error.\n");
705         exit(EXIT_FAILURE);
706     }
707     ret=posix_memalign((void**) &results,64,array_size*sizeof(int64_t));
708     if (ret)
709     {
710         fprintf(stderr, "Memory allocation error.\n");
711         exit(EXIT_FAILURE);
712     }
713
714     /* allocate the outer array and output arrays for band-join */
715     ret=posix_memalign((void**) &outer,64,outer_size*sizeof(int64_t));
716     if (ret)
717     {
718         fprintf(stderr, "Memory allocation error.\n");
719         exit(EXIT_FAILURE);
720     }
721     ret=posix_memalign((void**) &outer_results,64,result_size*sizeof(int64_t));
722     if (ret)
723     {
724         fprintf(stderr, "Memory allocation error.\n");
725         exit(EXIT_FAILURE);
726     }
727     ret=posix_memalign((void**) &inner_results,64,result_size*sizeof(int64_t));
728     if (ret)

```

```

729     {
730         fprintf(stderr, "Memory allocation error.\n");
731         exit(EXIT_FAILURE);
732     }
733
734
735     /* code to initialize data structures goes here so that it is not included in the timing
measurement */
736     init(data,queries,arraysize);
737     band_init(outer,outer_size);
738
739     #ifdef DEBUG
740         /* show the arrays */
741         printf("data: ");
742         for(int64_t i=0;i<arraysize;i++) printf("%ld ",data[i]);
743         printf("\n");
744         printf("queries: ");
745         for(int64_t i=0;i<arraysize;i++) printf("%ld ",queries[i]);
746         printf("\n");
747         printf("outer: ");
748         for(int64_t i=0;i<outer_size;i++) printf("%ld ",outer[i]);
749         printf("\n");
750     #endif
751
752
753     /* now measure... */
754
755     gettimeofday(&before,NULL);
756
757     /* the code that you want to measure goes here; make a function call */
758     printf("bulk_bin_search_start\n");
759     bulk_bin_search(data,arraysize,queries,arraysize,results, repeats);
760
761     gettimeofday(&after,NULL);
762     printf("Time in bulk_bin_search loop is %ld microseconds or %f microseconds per search\n",
(after.tv_sec-before.tv_sec)*1000000+(after.tv_usec-before.tv_usec), 1.0*((after.tv_sec-
before.tv_sec)*1000000+(after.tv_usec-before.tv_usec))/arraysize/repeats);
763
764
765
766     gettimeofday(&before,NULL);
767
768     /* the code that you want to measure goes here; make a function call */
769     printf("bulk_bin_search_4x_start\n");
770     bulk_bin_search_4x(data,arraysize,queries,arraysize,results, repeats);
771
772     gettimeofday(&after,NULL);
773     printf("Time in bulk_bin_search_4x loop is %ld microseconds or %f microseconds per search\n",
(after.tv_sec-before.tv_sec)*1000000+(after.tv_usec-before.tv_usec), 1.0*((after.tv_sec-

```

```

before.tv_sec)*1000000+(after.tv_usec-before.tv_usec))/arraysize/repeats);
774
775
776
777     gettimeofday(&before,NULL);
778
779     /* the code that you want to measure goes here; make a function call */
780     total_results=band_join(data, arraysize, outer, outer_size, inner_results, outer_results,
result_size, bound);
781
782     gettimeofday(&after,NULL);
783     printf("Band join result size is %ld with an average of %f matches per output
record\n",total_results, 1.0*total_results/(1.0+outer_results[total_results-1]));
784     printf("Time in band_join loop is %ld microseconds or %f microseconds per outer record\n",
(after.tv_sec-before.tv_sec)*1000000+(after.tv_usec-before.tv_usec), 1.0*((after.tv_sec-
before.tv_sec)*1000000+(after.tv_usec-before.tv_usec))/outer_size);
785
786 #ifdef DEBUG
787     /* show the band_join results */
788     printf("band_join results: ");
789     for(int64_t i=0;i<total_results;i++) printf("(%ld,%ld) ",outer_results[i],inner_results[i]);
790     printf("\n");
791
792 #endif
793
794
795     gettimeofday(&before,NULL);
796
797     /* the code that you want to measure goes here; make a function call */
798     total_results=band_join_simd(data, arraysize, outer, outer_size, inner_results, outer_results,
result_size, bound);
799
800     gettimeofday(&after,NULL);
801     printf("Band join (SIMD) result size is %ld with an average of %f matches per output
record\n",total_results, 1.0*total_results/(1.0+outer_results[total_results-1]));
802     printf("Time in band_join_simd loop is %ld microseconds or %f microseconds per outer
record\n", (after.tv_sec-before.tv_sec)*1000000+(after.tv_usec-before.tv_usec), 1.0*((after.tv_sec-
before.tv_sec)*1000000+(after.tv_usec-before.tv_usec))/outer_size);
803
804 #ifdef DEBUG
805     /* show the band_join results */
806     printf("band_join_simd results: ");
807     for(int64_t i=0;i<total_results;i++) printf("(%ld,%ld) ",outer_results[i],inner_results[i]);
808     printf("\n");
809 #endif
810
811
812 // FILE *csvFile;
813 // csvFile = fopen("band_join.csv", "w");

```

```
814 // fprintf(csvFile, "inner, outer\n");
815 // for(int64_t i=0;i<total_results1;i++){
816 // fprintf(csvFile, "%ld, %ld\n", inner_results[i],outer_results[i]);
817 // }
818 // fclose(csvFile);
819
820 // FILE *csvFile2;
821 // csvFile2 = fopen("band_join_SIMD.csv", "w");
822 // fprintf(csvFile2, "inner, outer\n");
823 // for(int64_t i=0;i<total_results2;i++){
824 // fprintf(csvFile2, "%ld, %ld\n", inner_results[i],outer_results[i]);
825 // }
826 // fclose(csvFile2);
827
828 // FILE *csvFile3;
829 // csvFile3 = fopen("data.csv", "w");
830 // fprintf(csvFile3, "inner, outer\n");
831 // for(int64_t i=0;i<arraysize;i++){
832 // fprintf(csvFile3, "%ld, %ld\n", data[i], outer[i]);
833 // }
834 // fclose(csvFile3);
835
836
837 }
838
839
```



```
1  /*
2   CSE 5242 Project 2, Fall 2023
3
4   See class project handout for more extensive documentation.
5
6   https://stackoverflow.com/questions/19068705/undefined-reference-when-calling-inline-function
7  */
8
9  #include <stdlib.h>
10 #include <stdio.h>
11 #include <stdint.h>
12 #include <unistd.h>
13 #include <string.h>
14 #include <sys/ioctl.h>
15 #include <sys/time.h>
16 #include <asm/unistd.h>
17 #include <immintrin.h>
18
19 /* uncomment out the following DEBUG line for debug info, for experiment comment the DEBUG line
20 */
21 // #define DEBUG
22
23 /* compare two int64_t values - for use with qsort */
24 static int compare(const void *p1, const void *p2)
25 {
26     int a,b;
27     a = *(int64_t *)p1;
28     b = *(int64_t *)p2;
29     if (a<b) return -1;
30     if (a==b) return 0;
31     return 1;
32 }
33
34 /* initialize searches and data - data is sorted and searches is a random permutation of data */
35 int init(int64_t* data, int64_t* searches, int count)
36 {
37     for(int64_t i=0; i<count; i++){
38         searches[i] = random();
39         data[i] = searches[i]+1;
40     }
41     qsort(data,count,sizeof(int64_t),compare);
42 }
43
44 /* initialize outer probes of band join */
45 int band_init(int64_t* outer, int64_t size)
```

```

46 {
47     for(int64_t i=0; i<size; i++){
48         outer[i] = random();
49     }
50 }
51
52 inline int64_t simple_binary_search(int64_t* data, int64_t size, int64_t target)
53 {
54     int64_t left=0;
55     int64_t right=size;
56     int64_t mid;
57
58     while(left<=right) {
59         mid = (left + right)/2; /* ignore possibility of overflow of left+right */
60         if (data[mid]==target) return mid;
61         if (data[mid]<target) left=mid+1;
62         else right = mid-1;
63     }
64     return -1; /* no match */
65 }
66
67 inline int64_t low_bin_search(int64_t* data, int64_t size, int64_t target)
68 {
69     /* this binary search variant
70      (a) does only one comparison in the inner loop
71      (b) doesn't require an exact match; instead it returns the index of the first key >= the search key.
72          That's good in a DB context where we might be doing a range search, and using binary search to
73          identify the first key in the range.
74      (c) If the search key is bigger than all keys, it returns size.
75     */
76     int64_t left=0;
77     int64_t right=size;
78     int64_t mid;
79
80     printf("Low_bin_search\n");
81     while(left<right) {
82         mid = (left + right)/2; /* ignore possibility of overflow of left+right */
83         if (data[mid]>=target)
84             right=mid;
85         else
86             left=mid+1;
87     }
88     return right;
89 }
90
91 //define ARDEBUB
92 inline int64_t low_bin_nb_arithmetic(int64_t* data, int64_t size, int64_t target)
93 {
94     /* this binary search variant

```

```

95     (a) does no comparisons in the inner loop by using multiplication and addition to convert control
dependencies
96     to data dependencies
97     (b) doesn't require an exact match; instead it returns the index of the first key >= the search key.
98     That's good in a DB context where we might be doing a range search, and using binary search to
99     identify the first key in the range.
100    (c) If the search key is bigger than all keys, it returns size.
101    */
102    int64_t left=0;
103    int64_t right=size;
104    int64_t mid;
105
106    #ifdef ARDEBUG
107        printf("low_bin_nb_arith\n");
108    #endif
109    //0^1 = 1 and 1^1 = 0
110    while(left<right) {
111        mid = (left + right) / 2; //get middle
112        int64_t flag = data[mid] >= target; //data[mid] >= target ? 1 : 0 (if target value is left side of the
middle index flag = 1 else 0)
113        left = flag * left + (flag^1) * (mid+1); // if flag == 1 assign previous left value (left index stays same)
else assign mid+1 value to divide the array into half
114        right = flag * mid + (flag^1) * right; // if flag == 1 assign mid value to divide array into half else assign
previous right value (right index stays same)
115        #ifdef ARDEBUG
116            printf("mid: %d left: %d right: %d\n", mid, left, right);
117        #endif
118    }
119    return right;
120 }
121
122 //define MASTDEBUG
123 inline int64_t low_bin_nb_mask(int64_t* data, int64_t size, int64_t target)
124 {
125     /* this binary search variant
126     (a) does no comparisons in the inner loop by using bit masking operations to convert control
dependencies
127     to data dependencies
128     (b) doesn't require an exact match; instead it returns the index of the first key >= the search key.
129     That's good in a DB context where we might be doing a range search, and using binary search to
130     identify the first key in the range.
131     (c) If the search key is bigger than all keys, it returns size.
132     */
133     int64_t left=0;
134     int64_t right=size;
135     int64_t mid;
136
137     #ifdef MASTDEBUG
138         printf("low_bin_nb_mask\n");

```

```

139 #endif
140 // ~(-1) = 0 and ~(0) = -1
141 //-1 & 3 = 3 and 0 & 3 = 0
142 //0 | 3 = 3
143 while(left<right) {
144     //if-else using bitwise source: https://stackoverflow.com/questions/3798601/conditional-using-bitwise-operator
145     mid = (left + right) / 2; //get middle
146     int64_t flag = (data[mid] >= target)-1; //data[mid] >= target ? 0 : -1 (if target value is left side of
middle index flag = 0 else =1)
147     left = (flag & (mid+1)) | (~flag & left); // If flag == -1 assign mid+1 to divide the array into half else
assign previous left value (left index stays same)
148     right = (flag & right) | (~flag & mid); // if flag == -1 assign previous right value (right index stay same)
else assing mid to divide the array into half
149     #ifdef MASTDEBUG
150         printf("mid: %d left: %d right: %d\n", mid, left, right);
151     #endif
152     //int64_t flag = (data[mid] >= target); //data[mid] >= target ? 1 : 0
153     // left = ((flag-1) & (mid+1)) | ((flag^1)-1 & left);
154     // right = ((flag-1) & right) | ((flag^1)-1 & mid);
155 }
156 return right;
157 }
158
159 //define XDEBUG
160 inline void low_bin_nb_4x(int64_t* data, int64_t size, int64_t* targets, int64_t* right)
161 {
162     /* this binary search variant
163     (a) does no comparisons in the inner loop by using bit masking operations instead
164     (b) doesn't require an exact match; instead it returns the index of the first key >= the search key.
165         That's good in a DB context where we might be doing a range search, and using binary search to
166         identify the first key in the range.
167     (c) If the search key is bigger than all keys, it returns size.
168     (d) does 4 searches at the same time in an interleaved fashion, so that an out-of-order processor
169         can make progress even when some instructions are still waiting for their operands to be ready.
170
171     Note that we're using the result array as the "right" elements in the search so no need for a return
statement
172     */
173     int64_t left[4]={0,0,0,0};
174     int64_t mid[4];
175     right[0]=right[1]=right[2]=right[3]=size;
176
177     #ifdef XDEBUG
178         printf("low_bin_nb_4x\n");
179     #endif
180     int64_t bit_size = 4;
181     int cnt = 0;
182     //outer loop to check termination condition like above (e.g left < right)

```

```

183 while (left[0]<right[0] || left[1]<right[1] || left[2]<right[2] || left[3]<right[3]) {
184     int i;
185     #ifdef XDEBUG
186         printf("cnt: %d\n", cnt);
187     #endif
188     //inner loop to perform 4 search concurrently
189     for (i = 0; i < bit_size; i++) {
190         //if-else using bitwise source: https://stackoverflow.com/questions/3798601/conditional-using-
        bitwise-operator
191         mid[i] = (left[i] + right[i]) / 2; //get middle
192         int64_t flag = (data[mid[i]] >= targets[i]) - 1; //data[mid] >= target ? 0 : -1 (if target value is left side
        of middle index flag = 0 else =1)
193         left[i] = (flag & (mid[i]+1)) | (~flag & left[i]); // If flag == -1 assign mid+1 to divide the array into half
        else assign previous left value (left index stays same)
194         right[i] = (flag & right[i]) | (~flag & mid[i]); // if flag == -1 assign previous right value (right index
        stay same) else assing mid to divide the array into half
195         #ifdef XDEBUG
196             printf("mid: %d left: %d right: %d\n", mid[i], left[i], right[i]);
197         #endif
198
199         //int64_t flag = (data[mid[i]] > targets[i]); //data[mid] >= target ? 1 : 0
200         // left[i] = ((flag-1) & (mid[i]+1)) | ((flag^1)-1 & left[i]);
201         // right[i] = ((flag-1) & right[i]) | ((flag^1)-1 & mid[i]);
202         //break;
203     }
204     cnt+=1;
205 }
206 }
207
208
209 /* The following union type is handy to output the contents of AVX512 data types */
210 union int8x4 {
211     __m256i a;
212     int64_t b[4];
213 };
214
215 void printavx(char* name, __m256i v) {
216     union int8x4 n;
217
218     n.a=v;
219     printf("Value in %s is [%ld %ld %ld %ld ]\n",name,n.b[0],n.b[1],n.b[2],n.b[3]);
220 }
221
222 /*
223  * Optinal for using AVX-512
224
225 union int8x8 {
226     __m512i a;
227     int64_t b[8];

```

```

228 };
229
230 void printavx512(char* name, __m512i v) {
231     union int8x4 n;
232
233     n.a=v;
234     printf("Value in %s is [%ld %ld %ld %ld %ld %ld %ld %ld]\n",name,n.b[0],n.b[1],n.b[2],n.b[3]);
235 }
236
237 */
238
239 //define SIMDDEBUG
240 inline void low_bin_nb_simd(int64_t* data, int64_t size, __m256i target, __m256i* result)
241 {
242     /* this binary search variant
243      (a) does no comparisons in the inner loop by using bit masking operations instead
244      (b) doesn't require an exact match; instead it returns the index of the first key >= the search key.
245          That's good in a DB context where we might be doing a range search, and using binary search to
246          identify the first key in the range.
247      (c) If the search key is bigger than all keys, it returns size.
248      (d) does 4 searches at the same time using AVX2 intrinsics
249
250      See https://www.intel.com/content/www/us/en/docs/cpp-compiler/developer-guide-reference/2021-
251      8/intrinsics-for-avx2.html
252      for documentation of the AVX512 intrinsics
253
254      Note that we're using the result array as the "right" elements in the search, and that searchkey is
255      being passed
256      as an __m256i value rather than via a pointer.
257      */
258
259     __m256i aleft = _mm256_set1_epi64x(0);
260     __m256i aright = _mm256_set1_epi64x(size);
261     __m256i amid;
262
263     __m256i ones = _mm256_set1_epi64x(1);
264     __m256i amask;
265     __m256i datavec;
266
267 #ifdef SIMDDEBUG
268     printf("low_bin_nb_simd\n");
269 #endif
270
271     //AVX Intrinsic Guide: https://www.intel.com/content/www/us/en/docs/intrinsics-guide/index.html
272     __m256i cmp = _mm256_cmpgts_epi64(aright, aleft);
273     //Check if all elements of vector is 0 using testz: https://stackoverflow.com/questions/32072169/could-i-compare-to-zero-register-in-avx-correctly + https://stackoverflow.com/a/32120039
274     //Check all elements in cmp are 0 (if value of cmp is 0 it means left > right else -1)
275     while(_mm256_testz_si256(cmp, cmp) == 0){
276         amid = _mm256_srli_epi64(_mm256_add_epi64(aleft, aright), 1); //get middle

```

```

274   _mm256i amid_one = _mm256_add_epi64(amid, ones); //get middle+1 vector
275   //Extract matching values of data using AVX index vector using i64gather:
https://stackoverflow.com/questions/51128005/what-do-you-do-without-fast-gather-and-scatter-in-
avx2-instructions
276   datavec = _mm256_i64gather_epi64((long long*)data, amid, sizeof(int64_t));
277
278   // inlining failed error for >= operation
279   //__mmask8 flag = _mm256_cmpgqe_epi64_mask(datavec, target);
280   //Combine > operation and = operation to have >= operation: https://www.splashlearn.com/math-
vocabulary/greater-than-or-equal-to
281   _mm256i cmpgt = _mm256_cmpgt_epi64(datavec, target);
282   _mm256i cmpeq = _mm256_cmpeq_epi64(datavec, target);
283   _mm256i flag = _mm256_or_si256(cmpgt, cmpeq);
284
285   //if-else using bitwise source: https://stackoverflow.com/questions/3798601/conditional-using-
bitwise-operator
286   aleft = _mm256_or_si256(_mm256_and_si256(flag, aleft), _mm256_andnot_si256(flag, amid_one)); // if
falg == -1 assign left value (left index stays same) else assing mid+1 to divide the array into half.
287   aright = _mm256_or_si256(_mm256_and_si256(flag, amid), _mm256_andnot_si256(flag, aright)); // if
flag == -1 assign mid value to divide array into half else assign previous right value (rihgt index stays
same)
288
289   #ifdef SIMDDEBUG
290       printavx("amid", amid);
291       printavx("flag", flag);
292       printavx("aright", aright);
293       printavx("aleft", aleft);
294       printf("\n\n");
295   #endif
296   //break;
297   cmp = _mm256_cmpgt_epi64(aright, aleft);
298   }
299   //aright = _mm256_sub_epi64(aright, ones);
300   *result = aright;
301   }
302
303   void bulk_bin_search(int64_t* data, int64_t size, int64_t* searchkeys, int64_t numsearches, int64_t*
results, int repeats)
304   {
305       for(int j=0; j<repeats; j++) {
306           /* Function to test a large number of binary searches
307
308               we might need repeats>1 to make sure the events we're measuring are not dominated by various
overheads, particularly for small values of size and/or numsearches
309
310               we assume that we want exactly "size" searches, where "size" is the length if the searchkeys array
311               */
312           for(int64_t i=0; i<numsearches; i++) {
313               #ifdef DEBUG

```

```

315     printf("Searching for %ld...\n",searchkeys[i]);
316 #endif
317
318     // Uncomment one of the following to measure it
319     //results[i] = low_bin_search(data,size,searchkeys[i]);
320     //results[i] = low_bin_nb_arithmetic(data,size,searchkeys[i]);
321     results[i] = low_bin_nb_mask(data,size,searchkeys[i]);
322
323 #ifdef DEBUG
324     printf("Result is %ld\n",results[i]);
325 #endif
326 }
327 }
328 }
329
330 void bulk_bin_search_4x(int64_t* data, int64_t size, int64_t* searchkeys, int64_t numsearches, int64_t*
results, int repeats)
331 {
332     register __m256i searchkey_4x;
333
334     for(int j=0; j<repeats; j++) {
335         /* Function to test a large number of binary searches using one of the 8x routines
336
337         we might need repeats>1 to make sure the events we're measuring are not dominated by various
338         overheads, particularly for small values of size and/or numsearches
339
340         we assume that we want exactly "size" searches, where "size" is the length if the searchkeys array
341         */
342         int64_t extras = numsearches % 4;
343         for(int64_t i=0; i<numsearches-extras; i+=4) {
344             #ifdef DEBUG
345                 printf("Searching for %ld %ld %ld %ld ...\\n",
346                     searchkeys[i],searchkeys[i+1],searchkeys[i+2],searchkeys[i+3]);
347             #endif
348
349             // Uncomment one of the following depending on which routine you want to profile
350
351             // Algorithm A
352             //low_bin_nb_4x(data,size,&searchkeys[i],&results[i]);
353
354             // Algorithm B
355             searchkey_4x = _mm256_loadu_si256((__m256i *)&searchkeys[i]);
356             low_bin_nb_simd(data,size,searchkey_4x,(__m256i *)&results[i]);
357
358             #ifdef DEBUG
359                 printf("Result is %ld %ld %ld %ld ...\\n",
360                     results[i],results[i+1],results[i+2],results[i+3]);
361             #endif
362         }

```



```

363  /* a little bit more work if numsearches is not a multiple of 8 */
364  for(int64_t i=numsearches-extras;i<numsearches; i++) {
365
366      results[i] = low_bin_nb_mask(data,size,searchkeys[i]);
367
368  }
369
370  }
371  }
372
373
374  int64_t band_join(int64_t* inner, int64_t inner_size, int64_t* outer, int64_t outer_size, int64_t*
inner_results, int64_t* outer_results, int64_t result_size, int64_t bound)
375  {
376      /* In a band join we want matches within a range of values. If p is the probe value from the outer
table, then all
377      records in the inner table with a key in the range [p-bound,p+bound] inclusive should be part of
the result.
378
379      Results are returned via two arrays. outer_results stores the index of the outer table row that
matches, and
380      inner_results stores the index of the inner table row that matches. result_size tells you the size of
the
381      output array that has been allocated. You should make sure that you don't exceed this size. If there
are
382      more results than can fit in the result arrays, then return early with just a prefix of the results in the
result
383      arrays. The return value of the function should be the number of output results.
384
385      */
386      //Declaring arrays to store left and right index of the range.
387      int64_t *leftIndexArray = malloc(outer_size * sizeof(int64_t));
388      int64_t *rightIndexArray = malloc(outer_size * sizeof(int64_t));
389      if (!leftIndexArray || !rightIndexArray) {
390          free(leftIndexArray);
391          free(rightIndexArray);
392          // Memory allocation failed, so we return -1
393          return -1;
394      }
395      //Calculating below values to find the multiplier and remainder.
396      int multiplier = outer_size/4;
397      int remainder = outer_size%4;
398
399      int temp_1=0;
400
401      while (temp_1<multiplier){
402          // printf("Inside multiplier with i->%d\n",temp_1);
403          //using the below arrays to store the lower bound values temporarily after using low_bin_nb_mask
function.

```

```

404     int64_t lower_limit[4];
405     int64_t upper_limit[4];
406
407     for (int i = 0; i < 4; i++) {
408         lower_limit[i] = outer[4 * temp_1 + i] - bound;
409         upper_limit[i] = outer[4 * temp_1 + i] + bound;
410     }
411     int64_t left_temp_Array[4];
412     int64_t right_temp_Array[4];
413
414     low_bin_nb_4x(inner, inner_size, lower_limit, left_temp_Array);
415     low_bin_nb_4x(inner, inner_size, upper_limit, right_temp_Array);
416
417     for (int i = 0; i < 4; i++) {
418         leftIndexArray[4 * temp_1 + i] = left_temp_Array[i];
419         rightIndexArray[4 * temp_1 + i] = right_temp_Array[i];
420     }
421     // printf("Ending multiplier with i->%d\n", temp_1);
422     temp_1 += 1;
423
424 }
425
426 for (int i = 0; i < remainder; i++) {
427     // printf("Entering remainder with i->%d\n", i);
428     // printf("lower-> %d\n", outer[i] - bound);
429     // printf("upper-> %d\n", outer[i] + bound);
430     leftIndexArray[4 * temp_1 + i] = low_bin_nb_mask(inner, inner_size, outer[4 * temp_1 + i] - bound); //array
containing left index in the range.
431     rightIndexArray[4 * temp_1 + i] = low_bin_nb_mask(inner, inner_size, outer[4 * temp_1 + i] + bound); //array
containing right index of the range.
432     // printf("leftindex->%d \t\n rightindex->%d\t\n", leftIndexArray[i], rightIndexArray[i]);
433     // printf("Exiting reminder loop with i->%d\n", i);
434 } printf("\n");
435
436
437
438 /* ONLY WITH low_bin_nb_mask
439 for (int i = 0; i < outer_size; i++) {
440     // printf("lower-> %d\n", outer[i] - bound);
441     // printf("upper-> %d\n", outer[i] + bound);
442     leftIndexArray[i] = low_bin_nb_mask(inner, inner_size, outer[i] - bound); //array containing left index in
the range.
443     rightIndexArray[i] = low_bin_nb_mask(inner, inner_size, outer[i] + bound); //array containing right index
of the range.
444     // printf("leftindex->%d \t\n rightindex->%d\t\n", leftIndexArray[i], rightIndexArray[i]);
445 } printf("\n");
446 */
447 // for (int i = 0; i < outer_size; i++) {
448 // printf("leftIndexArray %d->%d and rightIndexArray %d -

```

```

>%d\n",i,leftIndexArray[i],i,rightIndexArray[i]);
449 // }
450
451 int temp=0;
452
453 for(int i = 0; i<outer_size;i++){
454     // printf("Entering inside the merging part with i->%d\n",i);
455     int left_index = leftIndexArray[i];
456     int right_index = rightIndexArray[i];
457     //if left index and right index are same then we skip the iteration because there are no elements.
458     if (left_index==right_index){continue;}
459
460     while ((temp<result_size) && (left_index<right_index)){
461         // printf("Inside while for outer element index %d\n",i);
462         outer_results[temp]=i;
463         inner_results[temp]=left_index;
464         // printf("adding values (%d,%d) in outer and inner result\n",outer_results[i],inner_results[i]);
465         left_index+=1;
466         temp+=1;
467     }
468     //break if we reach the result size.
469     if (temp==result_size){break;}
470     // printf("Exiting merging part with i->%d\n",i);
471
472 }
473
474
475 // for (int i = 0; i<temp;i++){
476 //     printf("(OuterResult,InnerResult)->(%ld,%ld)\n",outer_results[i],inner_results[i]);
477 // }
478
479 free(leftIndexArray);
480 free(rightIndexArray);
481
482 return temp;
483
484 }
485
486 int64_t band_join_simd(int64_t* inner, int64_t inner_size, int64_t* outer, int64_t outer_size, int64_t*
inner_results, int64_t* outer_results, int64_t result_size, int64_t bound)
487 {
488     /* In a band join we want matches within a range of values. If p is the probe value from the outer
table, then all
489     records in the inner table with a key in the range [p-bound,p+bound] inclusive should be part of
the result.
490
491     Results are returned via two arrays. outer_results stores the index of the outer table row that
matches, and
492     inner_results stores the index of the inner table row that matches. result_size tells you the size of

```

the

```
493     output array that has been allocated. You should make sure that you don't exceed this size. If there
are
494     more results than can fit in the result arrays, then return early with just a prefix of the results in the
result
495     arrays. The return value of the function should be the number of output results.
496
497     To do the binary search, you could use the low_bin_nb_simd you just implemented to search for the
lower bounds in parallel
498
499     Once you've found the lower bounds, do the following for each of the 4 search keys in turn:
500     scan along the sorted inner array, generating outputs for each match, and making sure not to
exceed the output array bounds.
501
502     This inner scanning code does not have to use SIMD.
503     */
504
505
506
507 //Using similiar approach from the above function.
508
509 // Create and allocate space to store left and right result indexes
510 int64_t *leftIndexArray = malloc(outer_size * sizeof(int64_t));
511 int64_t *rightIndexArray = malloc(outer_size * sizeof(int64_t));
512 if (!leftIndexArray || !rightIndexArray) {
513     free(leftIndexArray);
514     free(rightIndexArray);
515     return -1;
516 }
517
518 // Calculate multiplier and remainder to find number of simd function instances and individual mask
function call instances
519 int multiplier = outer_size/4;
520 int remainder = outer_size%4;
521
522 int temp_1=0;
523 while (temp_1<multiplier){
524     // compute lower and upper bound vectors using AVX functions
525     __m256i outer_vec = _mm256_loadu_si256((__m256i*)&outer[4*temp_1]);
526     __m256i lower_bound_vec = _mm256_sub_epi64(outer_vec, _mm256_set1_epi64x(bound));
527     __m256i upper_bound_vec = _mm256_add_epi64(outer_vec, _mm256_set1_epi64x(bound));
528
529     // Create temporary indexes to store each iteration's results
530     int64_t left_indexes[4];
531     int64_t right_indexes[4];
532
533     // Call binary search function using simd
534     low_bin_nb_simd(inner,inner_size,lower_bound_vec,(__m256i*) left_indexes);
535     low_bin_nb_simd(inner,inner_size,upper_bound_vec,(__m256i*) right_indexes);
```

```

536
537 // Store results in temporary arrays
538 for (int i =0;i<4;i++){
539     leftIndexArray[4*temp_1 + i]=left_indexes[i];
540     rightIndexArray[4*temp_1 + i]=right_indexes[i];
541 }
542 temp_1+=1;
543
544 }
545
546 // Store indexes of remaining cases into the result arrays
547 for(int i = 0; i<remainder;i++){
548     leftIndexArray[4*temp_1 + i]=low_bin_nb_mask(inner, inner_size, outer[4*temp_1 + i]-bound);//array
containing left index in the range.
549     rightIndexArray[4*temp_1 + i]=low_bin_nb_mask(inner,inner_size,outer[4*temp_1 + i]+bound);//array
containing right index of the range.
550 }
551
552 // temp will be the final result size
553 int temp=0;
554
555 for(int i = 0; i<outer_size;i++){
556     int left_index = leftIndexArray[i];
557     int right_index = rightIndexArray[i];
558
559     if (left_index==right_index){continue;}
560
561     // Store indexes into the result format
562     while ((temp<result_size) && (left_index<right_index)){
563         outer_results[temp]=i;
564         inner_results[temp]=left_index;
565         left_index+=1;
566         temp+=1;
567     }
568 }
569
570 if (temp==result_size){break;}
571
572 }
573
574 // Clear index array space
575 free(leftIndexArray);
576 free(rightIndexArray);
577
578 // Return result size
579 return temp;
580
581
582 /*

```

```

583 INITIAL CODE WHICH WAS GIVING INCONSISTENT RESULT. PLEASE SKIP TO END OF COMMENT FOR
THE FINAL CODE.
584 int size = 0;
585
586 for(int i=0; (i<outer_size) && (size<result_size); i+=4){
587     __m256i outer_vec = _mm256_loadu_si256((__m256i*)&outer[i]);
588     __m256i lower_bound_vec = _mm256_sub_epi64(outer_vec, _mm256_set1_epi64x(bound));
589     __m256i upper_bound_vec = _mm256_add_epi64(outer_vec, _mm256_set1_epi64x(bound));
590
591     int64_t left_indexes[4];
592     int64_t right_indexes[4];
593
594     int remainder = outer_size - i;
595     if (remainder >= 4) {
596
597         __m256i left_result[1];
598         __m256i right_result[1];
599
600         low_bin_nb_simd(inner,inner_size,lower_bound_vec,left_result);
601         low_bin_nb_simd(inner,inner_size,upper_bound_vec,right_result);
602
603         _mm256_store_si256((__m256i*)left_indexes, left_result[0]);
604         _mm256_store_si256((__m256i*)right_indexes, right_result[0]);
605
606         for(int j = 0; j < 4; j++) {
607             while((size < result_size) && (left_indexes[j] < right_indexes[j])) {
608                 outer_results[size] = i+j;
609                 inner_results[size] = left_indexes[j];
610                 left_indexes[j]++;
611                 size++;
612             }
613         }
614     } else {
615
616         if (remainder >= 1) {
617             int64_t lower_bound = _mm256_extract_epi64(lower_bound_vec, 0);
618             int64_t upper_bound = _mm256_extract_epi64(upper_bound_vec, 0);
619
620             left_indexes[0] = low_bin_nb_mask(inner,inner_size,lower_bound);
621             right_indexes[0] = low_bin_nb_mask(inner,inner_size,upper_bound);
622         }
623
624         if (remainder >= 2) {
625             int64_t lower_bound = _mm256_extract_epi64(lower_bound_vec, 1);
626             int64_t upper_bound = _mm256_extract_epi64(upper_bound_vec, 1);
627
628             left_indexes[1] = low_bin_nb_mask(inner,inner_size,lower_bound);
629             right_indexes[1] = low_bin_nb_mask(inner,inner_size,upper_bound);
630

```

```

631     }
632
633     if (remainder == 3) {
634         int64_t lower_bound = _mm256_extract_epi64(lower_bound_vec, 2);
635         int64_t upper_bound = _mm256_extract_epi64(upper_bound_vec, 2);
636
637         left_indexes[2] = low_bin_nb_mask(inner, inner_size, lower_bound);
638         right_indexes[2] = low_bin_nb_mask(inner, inner_size, upper_bound);
639     }
640
641     for(int j = 0; j < remainder; j++) {
642         while((size < result_size) && (left_indexes[j] < right_indexes[j])) {
643             outer_results[size] = i+j;
644             inner_results[size] = left_indexes[j];
645             left_indexes[j]++;
646             size++;
647         }
648     }
649
650 }
651
652 }
653 return size;
654 */
655
656 }
657
658
659

```

▼ makefile


 Download

```

1  all: db5242
2
3  db5242: db5242.c
4      gcc -O3 -mavx2 -o db5242 db5242.c
5


```

▼ Project2 Document.pdf

 Download

Your browser does not support PDF previews. You can [download the file instead.](#)

▼ Project2 wo grammarly (3).docx

 Download

1

Large file hidden. You can download it using the button above.


```
1  # C-SIMD-AVX2-BS
2  C-Single Instructions Multiple Data (SIMD) using AVX2 to implement Binary Search
3
4  <h3>How to compile code</h3>
5
6  To use gcc to compile:
7
8  ```
9  gcc -O3 -mavx2 -o db5242 db5242.c
10 ```
11
12 To use makefile to compile
13
14 ```
15 make
16 ```
17
18 <h3>Ho to run compiled code</h3>
19
20 After compile code, db5242 class file should be created.
21
22 To run code:
23
24 ```
25 db5242 N X Y Z R
26 ```
27
28 where
29
30     N=size of array
31
32     X=Size of outer array
33
34     Y=Size of band join result
35
36     Z=bound
37
38     R=number of repeation for binary search
39
40
```